

Online Display Advertising Markets: A Literature Review and Future Directions

Hana Choi* Carl Mela[†] Santiago Balseiro[‡] Adam Leary[§]

May 10, 2019

Abstract

This paper summarizes the display advertising literature, organizing the content by the agents in the display advertising ecosystem, and proposes new research directions. In doing so, we take an interdisciplinary view, drawing connections among diverse streams of theoretical and empirical research in information systems, marketing, economics, operations, and computer science. By providing an integrated view of the display advertising ecosystem, we hope to bring attention to the outstanding research opportunities in this economically consequential and rapidly growing market.

Keywords: Display Advertising, Literature Review, Real-Time Buying, Programmatic Buying, Real-Time Bidding

JEL Classification Codes: M3, L1

*Hana Choi is a Ph.D. student at the Fuqua School of Business, Duke University (email: hana.choi@duke.edu, phone: 734-834-0699).

[†]Carl F. Mela is the T. Austin Finch Foundation Professor of Business Administration at the Fuqua School of Business, Duke University (email: mela@duke.edu, phone: 919-660-7767, web: faculty.fuqua.duke.edu/~mela).

[‡]Santiago R. Balseiro is Associate Professor of Business Administration at the Graduate School of Business, Columbia University (email: srb2155@columbia.edu, web: www.columbia.edu/~srb2155).

[§]Adam Leary is VP, Data Science Services at CBS Interactive (email: adam.leary@cbsinteractive.com, phone: 415-344-2645).

The authors thank Garrett Johnson for helpful comments and suggestions.

1 Introduction

According to the Interactive Advertising Bureau (IAB herein) report, display advertising (ad herein) totaled \$39.4 billion in the US in 2017, up from the \$31.2 billion reported in FY 2016.¹ This double-digit growth rate (26%) is fueled by the upswing in mobile browsing, social media activities, video ad formats, and the developments in targeting technology. Digital display ad spending is forecasted to surpass even search ad spending and will continue its rapid ascent in share and significance.²

Given the substantial economic importance of the display ad market, the objective of this paper is to summarize the existing literature from the perspective of each player involved in the display ad ecosystem. In doing so, we take an interdisciplinary view spanning diverse streams of theoretical and empirical research in information systems, marketing, economics, operations, and computer science. We focus on literature that explicitly emphasizes operations of the display ad market as opposed to a more general discussion on related topics, including the large body of literature on advertising, sponsored search ad,³ auctions, mechanism design, and online algorithms.

A handful of papers survey various aspects of online ad in general, covering search, display, classified, and other forms of ads. Ha (2008) reviews online ad research published specifically in six advertising journals.⁴ Evans (2008, 2009) discusses the evolution of internet ad and provides industry perspectives,⁵ while Bucklin Hoban (2017) focus on marketing models developed to improve decision-making in internet advertising. Due to the broad scope of these papers, the distinguishing characteristics of the display ad market (such as the issues arising from the coexistence of guaranteed and non-guaranteed selling channels) and its ecosystem are precluded.

Focusing on the display ad market, existing survey papers either discuss issues facing advertisers (demand side), publishers (supply side), or intermediaries. On the advertiser side, Tucker (2012a) and Goldfarb (2014) review the economic literature on targeting and privacy concerns, and highlight their trade-off in ad effectiveness. On the publisher side, Korula et al. (2016) cover a broad range

¹ https://www.iab.com/wp-content/uploads/2018/05/IAB-2017-Full-Year-Internet-Advertising-Revenue-Report.REV_.pdf

² Search advertising revenue totaled \$40.6 billion in 2017, a little above display advertising with a growth rate of 17%. Source: iab.com

³ Online Appendix highlights similarities and dissimilarities of online display advertising from sponsored search advertising.

⁴ Ha (2008) reviews articles published in *Journal of Advertising*, *Journal of Advertising Research*, *Journal of Current Issues and Research in Advertising*, *International Journal of Advertising*, *Journal of Marketing Communications*, and *Journal of Interactive Advertising*.

⁵ For a brief description of the history of online display advertising, please refer to Online Appendix.

of publishers’ decisions in selling display ad, and review studies on allocating impressions and designing contracts/auctions. On the intermediary side, Muthukrishnan (2009) considers issues faced by ad exchanges, which are online platforms that match advertisers with publishers. One exception is Wang et al. (2016), who discuss recent algorithms and challenges in computational advertising from both advertisers’ (bidding strategies in real-time buying) and publishers’ (pricing and auction mechanism) perspectives, focusing on the non-guaranteed display ad selling channel.

Our survey differs in that we provide a comprehensive view of the display ad ecosystem, including both guaranteed and non-guaranteed selling channels, and delineate the research progress gained from the eyes of all the players involved, including advertisers, publishers, and intermediaries. Furthermore, we take an interdisciplinary view and draw connections across disciplines. By providing an integrated view and bringing attention to the outstanding research opportunities, we hope to provide additional motivation for subsequent studies in this rapidly growing market.

The remainder of the paper is organized as follows. Section 2 overviews the ecosystem and highlights key unique features of the display ad market. Next, we organize existing literature by advertisers (Section 3), publishers (Section 4), and intermediaries (Section 5). In each sub-section, we address key managerial decisions and discuss related research issues and opportunities. Section 6 characterizes some contextual drivers of display advertising, namely mobile and social media, and Section 7 concludes.

2 The Display Advertising Ecosystem

This section enumerates the players, the selling channels, and the key characteristics of the display ad market. Players’ objectives, roles, and available information are briefly discussed to place in context the research topics discussed in subsequent sections.

2.1 Players

The display ad market is a two-sided market: on one side, advertisers procure ad impressions on publishers’ websites to reach potential consumers; and on the other side, publishers with consumers’ impressions purvey ad inventory to advertisers with the highest valuations for those impressions. In between, intermediaries facilitate the match between advertisers and publishers by managing data and providing optimization tools and algorithms for serving ads.

2.2 Selling Channels

Ad exposures in display markets are sold via guaranteed and non-guaranteed selling channels. The guaranteed selling channel (so-called direct sale) involves the advance sale of a number of impressions at a fixed price. Non-guaranteed sales occur in real time via ad exchanges (so-called Real-Time Buying, RTB herein).⁶ The non-guaranteed selling channel is estimated to be 35% ($\approx 82.5\% * 42\%$) of the total display ad revenue in 2018.⁷

The main characteristics distinguishing the guaranteed selling channel from RTB are (i) the pricing mechanism (fixed price in the guaranteed vs. auction in RTB), (ii) the information available about the impressions (thus targeting ability), and (iii) the players involved (especially the level of intermediation). These distinguishing characteristics are described for each selling channel next.

2.2.1 Guaranteed Selling Channel

The guaranteed selling channel involves the buying and selling of bundles of impressions through a guaranteed contract, prior to users' arrivals on the publisher's site. In the guaranteed selling channel, an advertiser and a publisher negotiate a fixed price pertaining to when, where, and how the ads will be displayed. These contractual arrangements guarantee the number of impressions to be delivered satisfying certain targeting criteria, at a negotiated fixed price (cost-per-mille, CPM), during a specified time period (e.g., 1M impressions to female users in SF during July for \$5000).⁸

Advertisers often work with ad agencies to facilitate media planning/buying and the negotiation process. On the publisher side, there often is an in-house sales team who negotiates the terms and prices of the contract and manages the relationship with the advertisers. Many aspects of the contracting process are being automated with programmatic buying technologies.⁹ The so-called *programmatic direct* (or programmatic guaranteed, automated guaranteed) enables publishers to avail their guaranteed deals to advertisers with smaller budgets who otherwise would not meet the minimum ad spend required to allocate the limited resources of the sales team. The advent of *programmatic direct* thus decreases the role of sales teams and lowers advertisers' search costs

⁶ We use "non-guaranteed selling channel" and "RTB" interchangeably in this article.

⁷ <https://www.emarketer.com/content/more-than-80-of-digital-display-ads-will-be-bought-programmatically-in-2018>

⁸ Instead of guaranteeing a number of impressions to be delivered, contracts may alternatively be written to guarantee delivery dates or a share of impressions (Danaher et al. 2010) at a fixed price. Although less common, publishers may guarantee the number of clicks to be delivered based on cost-per-click (CPC).

⁹ The terminology "programmatic buying" should not be confused with non-guaranteed RTB. Programmatic buying consists of a wide range of technologies that *automate* the buying, selling, and matching of ad.

associated with finding prices from multiple publishers, which in turn would affect advertisers' optimal ad buying strategies across selling channels (§3.2.1).¹⁰

Because guaranteed contracts are drawn in advance, both advertisers and publishers rely on their expectations about future impressions and associated user characteristics (of those who will arrive during the planned campaign period) when making contractual decisions. Because targeting criteria are typically coarse, advertisers end up buying a “bundle” of impressions at a fixed price. For an advertiser, important questions are when/how should this guaranteed bundle of impressions be bought and which targeting criteria to use (§3.2). For a publisher, the questions are whether to sell impressions in advance as a bundle and how to price the guaranteed inventory (§4.1).

2.2.2 Non-Guaranteed Selling Channel

The non-guaranteed selling channel involves the buying and selling of single impressions in real time via programmatic buying. This channel is called RTB, because an advertiser's buying decision is made immediately after the user arrives on the website. The ad inventory is non-guaranteed, because an auction takes place for the incoming impression and the advertiser's ad is served to the end user only if the advertiser wins the auction. RTB enables publishers to monetize beyond the guaranteed contracted volume and advertisers to buy impressions based on user-specific (e.g., cookie) information. Questions advertisers face in this market include how to calculate the value of an impression and bid, how to leverage users' past behavioral data, and how to react to competitors' bidding strategies (§3.2.3). Publisher questions include how to allocate an incoming impression across advertisers with guaranteed contracts and the highest bidder in RTB (§4.3).

The ecosystem in RTB is more complex than the guaranteed selling channel, with many specialized intermediaries lying in between. These intermediaries provide the fundamental infrastructure required for selling, buying, and serving ads in real time, typically within milliseconds.¹¹ For example, *Ad exchanges* (e.g., DoubleClick Ad Exchange, OpenX, AppNexus Marketplace) provide centralized marketplaces by conducting an auction for each available impression in real time. Recognizing that programmatic buying is a difficult operational problem, *Demand side platforms* (DSPs herein; e.g., MediaMath, Rocket Fuel, Google Marketing Platform - Display & Video 360)

¹⁰ Depending on the publisher, negotiation is altogether precluded in *programmatic direct* to facilitate the automated contracting process. For example, the quoted price in Facebook's “Reach and Frequency Buying” is a take-it-or-leave-it fixed price.

¹¹ <http://data.iab.com/ecosystem.html>

help manage placement and bids for the advertisers. Similarly on the publisher side, *Supply side platforms* (SSPs herein; e.g., DoubleClick for Publishers, Rubicon Project for Sellers, PubMatic, MoPub) facilitate publishers' inventory management and yield optimization. Although little research exists to date on the role of these intermediaries, the welfare of advertisers and publishers may improve by intermediaries' enabling technologies that enhance targeting and allocation of impressions, although concerns exist regarding high commissions and incentive mis-alignment (§5.1, §5.2).

Finally, there are data intermediaries. Both advertisers and publishers strive to manage and leverage consumer data efficiently for targeting, but have access to different information. Advertisers, for example, can augment their own proprietary data (e.g., purchase patterns on the advertiser's website) with additional data bought from third parties (e.g., income level, job history, home ownership, monthly car payment) in order to build targeting audience profiles. Publishers seek to integrate their own proprietary data (e.g., user registration information and ad viewing patterns on the publisher's website) with third-party data to build user segments and to offer better targeting options to advertisers at potentially higher prices. To this end, *Data Management Platforms* (DMPs herein; e.g., Adobe Audience Manager, Lotame, SAS Data Management, Oracle DMP, Nielsen DMP, Salesforce DMP) provide an accessible interface to import data from multiple sources and to build audience segments from the integrated data, which are then fed into DSP/SSP toward optimizing downstream decisions. For both advertisers and publishers, important research questions pertain to the optimal level of targeting granularity (to buy or to sell) and to which additional information to buy from third parties and at what price (§3.3.3, §4.2). Asymmetry of information also raises questions regarding how much information to share and the role of intermediaries in resolving this asymmetry (§5.4).

3 Advertisers

Media planning and execution constitute the demand side of the display ad market. Advertisers (i) set objectives (§3.1) and (ii) buy a set of impressions toward maximizing the objectives given their constraints (§3.2). Throughout the ad campaign, marketers (iii) measure return on investment (ROI herein) and use it as a reference point in setting future objectives and buying strategies (§3.3).

3.1 Objective

Characterizations of display advertising have suggested two major goal types (e.g., Zhu Wilbur, 2011). One type involves building long-term brand equity. The other type is to generate short-term direct responses similar to coupons in print ads. As display ad objective is not directly observed by researchers in most datasets, integrated studies about what factors drive this twofold objective and the relative importance of goal types are scant. As display ad objective drives advertisers’ downstream decisions (such as budgeting and ad buying) and is antecedent to modeling, understanding this twofold objective is a pertinent area of future research.

There exists an array of key performance indicators (KPIs herein) used by advertisers to evaluate their campaign performance and to set (implementable) objective functions for optimization algorithms. The most common amongst these include impressions (Danaher et al. 2010), clicks (Chatterjee et al. 2003), visits (Dalessandro et al. 2015), and purchases (Manchanda et al. 2006). Consequent to the often incongruent twofold objective of increasing brand equity and enhancing direct response, the question arises of which metrics should be evaluated and optimized. For example, should advertisers contract to buy a guaranteed number of impressions versus clicks? Does the time spent on a page and the number of page views constitute good measures for advertisers focusing on branding? Should survey-based metrics such as brand awareness, attitude, recognition, or recall be collected pre- and/or post-campaign (Goldfarb Tucker 2011a; Bart et al. 2014)? If so, how can these data be leveraged in the optimization process? Are advertisers using the right metrics to drive profitability? Future research can further explore the relationship among advertisers’ display ad objectives, optimization processes, and various KPIs evaluated in practice.

3.2 Ad Buying

After the advertising objective is set, advertisers face the following ad buying decisions: (i) whether to buy guaranteed or non-guaranteed (§3.2.1), (ii) ad design and inventory characteristics (§3.2.2), (iii) user characteristics (targeting) (§3.2.3), and (iv) scheduling (frequency and timing) (§3.2.4). We discuss each decision in turn.

3.2.1 Selling Channels: Guaranteed vs. Non-Guaranteed

Most advertisers engage both selling channels by allocating a portion of their budget in the guaranteed selling channel in advance, and then spending the rest in RTB. Unfortunately, the existing

literature on display ad is largely silent on why/when one selling channel is/should be chosen over the other. Below we discuss some plausible explanations, which future research can explore.

Some reasons why advertisers may favor RTB include finer targeting and lower search costs. When buying in real time, advertisers can submit an appropriate bid based on user-specific information at the impression level (e.g., user’s past behavioral pattern). Further, procuring a guaranteed contract entails time-consuming negotiation processes with potentially multiple publishers. These search costs are significantly reduced in RTB, as advertisers can reach a large number of publishers through a centralized intermediary exchange market.

On the other hand, guaranteed contracts are often favored when advertisers have a long-standing relationship with the publisher that facilitates the customization of ad formats and price negotiations (e.g., quantity discounts, bundling). If advertisers are risk averse, paying a premium to buy a guaranteed inventory in advance helps mitigate uncertainty in either the auctions’ outcomes or the amount of impressions that will be available in RTB on specific dates. Lastly, advertisers highly concerned with ensuring brand safety will choose guaranteed contracts so their ads appear on high-quality, reputable websites.¹²

Athey et al. (2017) demonstrate that with consumers’ multi-homing across publishers and imperfect tracking technology, advertisers seeking broader reach would favor larger publishers to avoid inefficient duplication, which can create an incentive for contracting guaranteed deals.¹³ As a supply side explanation, Sayedi (2018) explores dynamic allocation (see §4.3), under which having both selling channels yields higher profits than selling in the guaranteed selling channel or RTB alone.

A large number of positive and normative research questions remain regarding dual selling channels. Positively speaking, an issue of interest is what the rationales are behind advertisers’ observed selling channel choices and their relative buying shares. For example, why would some advertisers invest more heavily in the guaranteed contracts and others more heavily in RTB? Why do dual selling channels co-exist in equilibrium? Normatively speaking, when would it be beneficial for the advertisers to choose one selling channel over the other, and what will be the optimal buying strategy?

¹² A 2017 survey by the Association of National Advertisers indicates that 78% of advertisers are concerned about brand safety issues in programmatic media buying, due to low transparency in working with external agencies. (<https://www.ana.net/content/show/id/47123>).

¹³ In their working paper version (Athey et al. 2014), Section 7 explores publishers offering guaranteed versus non-guaranteed deals.

3.2.2 Ad Design and Inventory Characteristics

In addition to the selling channels, advertisers make decisions regarding ad design and inventory characteristics.¹⁴ These involve choices about (i) websites (Danaher 2007; Danaher et al. 2010; James et al. 2018; Paulson et al. 2018) including traditional media (e.g., nytimes.com), blogs, and social media (e.g., Facebook) ; (ii) positions (Aksakalli 2012); (iii) sizes; (iv) creatives (Urban et al. 2013; Bruce et al. 2017); (v) formats including banner, text (e.g., Google AdSense), rich media (e.g., expanding creative), digital video, and native ads (Wojdyski Evans 2016); and (vi) devices including desktop, mobile, and tablets.

Past research has typically evaluated each ad design and inventory characteristic independently to ascertain its respective effects on traditional brand equity measures (e.g., awareness, recall, recognition, attitude) and/or direct response measures (e.g., CTR, purchase intent). However, relatively little effort has been apportioned toward assessing the inter-relationship among these various inputs or their efficacy with respect to cost (Johnson Lewis 2015). As ad design and inventory characteristics work jointly, optimization entails taking a combination of websites, positions, sizes, ad creatives, formats, and devices into account, and weighing the (incremental) benefits and costs of these options toward achieving the display ad goal given constraints (e.g., budget).

Furthermore, the choices of ad design and inventory characteristics are closely tied with the choices of selling channels. In RTB, standardization is necessary to ensure automation and seamless delivery within milliseconds. As a result, ad exchanges generally support inventories that conform to the IAB standard guideline, and customizability in terms of ad position, size, format, or content environment is limited.¹⁵ For example, highly customized ad positions and environments, such as homepage takeover (Sayedi et al. 2018) or conquest ads, are generally supported only via guaranteed contracts.¹⁶ Goldfarb Tucker (2014) provide evidence that ad recognition declines as the use of IAB standard formats rises, presumably because standard format ads attract less attention and are harder to distinguish from competitors' ads. This raises the question of how the level of standardization or the demand for customizability affect ad buying decisions and market outcomes.

From an implementation perspective, scalability, computational efficiency, and accuracy are key

¹⁴ See Ha (2008) and references therein for earlier research on advertisers' ad design and inventory decisions.

¹⁵ <https://www.iab.com/guidelines/>

¹⁶ Best Buy, for example, can *takeover* CNET's entire front page on Black Friday and display its creatives with unique design and architecture. Or Samsung can place *conquest* ads next to publisher's review article on Apple's newly launched iPhone, to reach those considering a competitor's product.

considerations in building optimization algorithms. An optimization algorithm taking both selling channels into account is imperative. For computational efficiency, advertisers might be making these decisions hierarchically in practice (e.g., allocating budget across selling channels first, then across websites, then devices and formats, etc.) and an easy-to-implement optimization algorithm might be considered upon these heuristics.

3.2.3 Target User Characteristics

Advertisers often target ads to users whose purchase might be imminent (generating direct responses) or who might potentially become valuable customers in the future (building brand equity or customer lifetime value). Central to this goal is the identification of targets conditioned upon the available information. In *demographic targeting*, users are segmented based on age, gender, and geography. *Contextual* advertising, on the other hand, targets based on users' interests; for example, by matching ad product category to the website content the user is browsing. In *behavioral targeting*, users are sought after based on their past behavioral patterns; for example, users who visited an advertiser's website in the past. We begin with targeting in the guaranteed setting, then move to the discussion of targeting in RTB.

Targeting Under Guaranteed Selling Channel For behavioral targeting, user profiles are generated based on users' topic interests (Ahmed et al. 2011; Trusov et al. 2016), conversion intents (Chung et al. 2010; Pandey et al. 2011; Aly et al. 2012), or brand affinity (Provost et al. 2009), and the preferred users are selected as the target audience. Contracting upon such preferred user segments several months in advance via guaranteed contracts, however, is often difficult as users' characteristics (e.g., topic interests) vary considerably over time. From a publisher standpoint, committing to deliver a certain number of impressions of users demanded by the advertiser may increase the risk of under-delivery, as the publisher may not have enough information to make good predictions on those users' future impressions. Although behavioral targeting based on the advertiser-side information (e.g., users identified as having high purchase intent based on previous visits to the advertiser's website) is not common, some publishers offer behavioral targeting based on the publisher-side information (e.g., users with frequent visits to the publisher's website). Presumably the publisher tries to price discriminate and increase profits by offering those user segments whose future impressions are well predicted by the publisher based on its own data, and

whose exposures are highly valued by the advertisers.

In terms of demographic and contextual targeting, advertisers often have good insights on which demographic groups and contents are compatible with the ad product or service. For example, the advertiser may want to target users living close to their brick-and-mortar store locations (Johnson et al. 2016a). The question advertisers often face is whether it is profitable to pay for an additional layer of targeting. Although not much research has studied advertisers’ user targeting decisions in the guaranteed selling channel, this is an eminent area for future research, especially in conjunction with RTB. For example, when targeting in guaranteed contracts is limited, advertisers might opt to buy inventory in RTB auctions where targeting is finer. The shift in bargaining power between advertisers and publishers will affect how much information (targeting) is provided and the price and demand for such targeting in the guaranteed selling channel.

Targeting Under Non-Guaranteed Selling Channel In RTB, advertisers have considerably more information and control over targeting. Each impression is sold separately in real time, meaning that the value of an impression can more readily be evaluated by matching the incoming user with the advertiser’s own data on users’ past behavior. In addition to the demographic and contextual targeting, advertisers greatly enjoy behavioral targeting in RTB.

Determining targeted optimal bids requires two components: (i) the advertiser’s own valuation for the targeted impression and (ii) distribution of the highest bids among competing advertisers. Aggregating targeted impressions into a campaign, the advertiser will calculate the optimal bids to maximize its overall utility (welfare) given the auction format. We first review the literature that address issues in *predicting* advertisers’ valuations and the distribution of others’ bids in real time, then move to the discussion on the optimal bidding strategy in consideration with the advertiser’s display ad goal and constraints.

1) **Real-Time Prediction on Advertiser Valuations**

Placing a bid entails predicting the advertiser’s value of an arriving impression in real time. Toward this end, prior research focuses upon using immediate ad responses (mainly CTR, conversion rates) as proxies for valuation and builds algorithms to minimize prediction errors. Some computational challenges in predicting direct ad responses in real time include that (i) the dimensionality of user attributes space is high and data is sparse across these dimensions, (ii) conversion events occur rarely, and (iii) cold-start situations with limited data history hinder reliable predictions. In

general, the literature has evolved to address these challenges and to lower computational costs while improving predictions. Linear models (Chen et al. 2009; Agarwal et al. 2010, 2014; McMahan et al. 2013; Chapelle et al. 2015) are widely used in practice as they are easy to implement and perform effectively even with large-scale datasets. As it is difficult to capture higher order interactions with linear models, hybrid methods are being developed, incorporating factorization (Menon et al. 2011; Oentaryo et al. 2014b; Juan et al. 2016) or decision trees (He et al. 2014), and deep learning methods are being applied to further explore latent patterns (Zhang et al. 2016b). The learnt user characteristics from user profiling (Zhang et al. 2014; Perlich et al. 2014) and features additionally inferred from social data (Bagherjeiran Parekh 2008; Liu Tang 2011; Goel Goldstein 2013) are also leveraged in improving ad response predictions.

An inherent presumption in much of the preceding literature is that valuations are proxied by immediate conversion probabilities. However, this practice suggests serving ads to consumers who are likely to convert even without ad exposures. A better bidding strategy would, for example, instead use *incremental CTR caused* by ad exposures as a proxy for advertiser valuations (Lewis Wong 2018). Branding is also relevant. Beyond maximizing the number of ad impressions, an ideal algorithm would maximize (incremental) expected profits over a long period of time by predicting customers’ lifetime value. Models of real-time prediction of advertiser valuations could be extended in this manner to better reflect true opportunities from the impression.

2) Real-Time Prediction on Distribution of the Highest Bids

The distribution of competing advertisers’ highest bids governs winning probabilities and expected payments. A bidder’s winning probability also determines the total expected number of impressions won and the *reach* metric for a given campaign. Challenges in predicting competing advertisers’ highest bids include that (i) impressions possess demographic, contextual, and behavioral attributes for which competing advertisers have heterogeneous valuations, and (ii) because competitors’ bids are not disclosed, the highest of the competing bids (and the reservation price set by the publisher) is only observed from the winning payment, and the lower bound of competing advertisers’ highest bids is inferred upon losing from the advertiser’s own bid. Toward addressing the first point, Cui et al. (2011) assume the winning bid distribution to be a mixture of log-normal distributions, where the mixture weights reflect different targeting features. The second point can be addressed by extending a censored regression model (Wu et al. 2015) or survival model (Wang Zhang 2016)

to correct for the bias in learning due to the difference in training data (observed, censored winning price history) and the testing data.

3) Optimal Bidding Strategy

Advertisers seek to determine their optimal bids given the auction mechanism faced. Until recently, ad exchanges typically conducted second-price, sealed-bid auction for each impression available, and much of the existing display advertising literature thus far studies this auction format (or a variant thereof).¹⁷ In theory, the weakly dominant strategy for an advertiser in a private value, second-price auction of a single object is bidding truthfully. However, bid calculation becomes more complex as additional practical constraints are considered: the advertiser (i) faces a budget constraint for a given campaign in repeated auctions (Balseiro et al. 2015; Balseiro Gur 2017), (ii) learns own and/or other’s true valuations over time (Iyer et al. 2014; Cai et al. 2017), (iii) sets a number of impressions to attain (Ghosh et al. 2009b; Choi Mela 2018), and (iv) sets pacing options so that the budget is spent smoothly over a specified time period (Lee et al. 2013; Yuan et al. 2013; Xu et al. 2015).¹⁸ With such considerations, the optimal bidding strategy is not necessarily truth telling. For example, Balseiro et al. (2015) show that the optimal bidding strategy for an advertiser facing a (binding) budget constraint is to shade values to account for the option value of future opportunities.

In designing optimal bidding algorithms, different assumptions have been made about the advertiser’s objective function and the bidding environment, including the ways predictions are made about valuations and the distribution of highest bids. One common approach predicts statistics of interest (e.g., CTR for ad responses, moments for the distribution of highest competing bids) as a preliminary step toward minimizing prediction errors, and then using these predictions as inputs in calculating the bids (Perlich et al. 2012; Zhang et al. 2014; Xu et al. 2016). However, when advertisers are highly uncertain about ad responses or the distribution of the highest bid, incorporating the prediction step into the entire bidding optimization task is desirable to enhance learning about

¹⁷ The first-price, sealed-bid auction format is becoming increasingly popular due primarily to the adoption of header bidding (see §5.3). Solving the advertiser bidding problem under a first-price auction is computationally more complex, as optimal bids deviate from advertiser valuations (e.g., bid shading), but is an especially topical consideration (<https://www.emarketer.com/content/understanding-programmatic-auction-pricing-is-a-major-priority-for-marketers>).

¹⁸ Spending the budget relatively smoothly over the campaign period (or over a given day or part of a day) is considered to help advertisers attain a wider range of audience. However, the commonly offered pacing option by the DSP, *uniform* pacing, may not be optimal as spending the budget absolutely evenly does not take into account the changes in the traffic, the number of competing advertisers, and the value of impressions over the course of time.

the overall utility from a campaign (Hummel McAfee 2015; Chapelle 2015; Ren et al. 2016).

Because training the model is computationally expensive, training has traditionally been done offline to find the optimal bidding strategy, which is then applied to the real-time data. As the campaign progresses, advertisers receive additional data on winning outcomes with corresponding payments and user responses. In order to incorporate such new information or to dynamically adapt to the changes in the competition environment, while moderating computational cost, research considers shortening the successive offline optimization cycles by applying simple algorithms that can myopically re-optimize (Lang et al. 2012) or by tuning the bidding parameters dynamically according to the current performance feedback (Cai et al. 2017).

Developing an adaptive and better performing real-time bidding algorithm is an ongoing topic of interests (e.g., Tunuguntla Hoban 2018). In this regard, the participation and transaction cost incurred in RTB are yet to be explored. At the selling channel level, advertisers face entry costs in terms of understanding the complexity in RTB and setting up an appropriate bidding process. At the auction level, there are transaction costs such as monetary fees charged by the intermediaries and/or cognitive effort costs in bidding. Although we are aware of no research quantifying the participation and transaction costs in RTB, these costs will considerably affect advertisers’ participation decisions, competitive landscape, and resulting auction outcomes.

3.2.4 Scheduling

We next consider advertisers’ ad scheduling decisions; that is, when to display ads and how often. Finding optimal frequency is important not only to attain the highest ad effectiveness (Johnson et al. 2016a) but also to optimally allocate limited resources. Timing is important due to the long-term carry-over or spacing effect of ads (Braun Moe 2013; Sahni 2015; Sahni et al. 2019).

In the guaranteed selling channel, publishers mostly control the ad delivery schedule and commonly offered decision variables are frequency capping and pacing options.¹⁹ An implementable algorithm that can choose these common decision variables (e.g., how to set frequency capping, does uniform pacing result in better exposure timing) toward achieving optimal scheduling will be valuable.

In RTB, where each impression can be bought separately based on the user-specific information,

¹⁹ Frequency capping restricts the number of times a particular campaign is displayed to a user, and pacing is the rate at which the budget is spent during the campaign. Uniform pacing involves spending the budget evenly, whereas no pacing refers to spending as fast as possible.

advertisers have better control over frequency and timing at the individual level, conditioned upon users’ browsing behaviors. To this end, one possible direction for future research is building an optimal bidding algorithm that takes the effect of frequency and timing into account (e.g., carry-over effect of ad) in evaluating own valuations of an impression and calculating the bid. Having a greater control over individual-level frequency and timing might be an additional rational for advertisers to allocate budget toward RTB. As winning is probabilistic in RTB and the bid landscape changes rapidly, ad scheduling is an interesting dynamic control problem.

3.3 Measuring Value of Ad Spend

Below we discuss challenges in measuring the actual post-hoc value of ad spend to the advertiser (i.e., causal effect as opposed to the advertiser’s ex-ante prediction for the ad valuation) and suggest some directions for future research.

3.3.1 Attribution

Assessing and quantifying the effectiveness of each ad is remarkably difficult, as the consumer is exposed to an advertiser multiple times via various channels, and as these “multi-touches” jointly influence consumer behavior. For example, display and TV ads simultaneously affect online and offline sales. How advertisers attribute the increase in desired outcome to each ad exposure affects how advertisers would update budget allocation across media channels and also the buying decisions in the display ad market (Jordan et al. 2011; Shao Li 2011; Abhishek et al. 2012; Dalessandro et al. 2012; Berman 2015; Li Kannan 2014; Xu et al. 2014; Kireyev et al. 2016). With the advancement in tracking technology, richer datasets become available that connect ad exposures not only across communication channels (display, search, TV) but also across devices (desktop, mobile, tablets). Individual-level online activities are further being matched to offline behaviors, for example, via reward cards. Ongoing progress is being made to capture partial aspects of these cross-channel effects (e.g., effect of display ad on online and offline sales in Lewis Reiley 2014; effect of search and display ads on online and offline purchases in Abraham 2008; effect of display ads on consumer searches in Papadimitriou et al. 2011, Lewis Nguyen 2015, Ghose Todri-Adamopoulos 2016; effect of mobile and desktop exposures on mobile and desktop conversions in Ghose et al. 2013) and future research is warranted encompassing the entire customer journey (i.e., effects of multi-channel exposures on multi-channel behaviors) to better understand attribution and the resulting

implications on ad buying strategies.

3.3.2 Econometric Issues and Experimental Design

Putting aside the attribution and cross-channel issues, measuring the causal effect in general poses several challenges, including selection biases arising from targeting and browsing intensity (see Lewis et al. 2015 and references therein, Hoban Arora 2018). Observational variation, even with the most sophisticated estimation methods, can easily fail to yield reliable measurement of ROI (Gordon et al. 2016). To overcome the econometric issues and to better understand the value of ad spend, both advertisers and researchers are turning attention to randomized field experiments. The key question is how to design experiments to obtain unbiased estimates with statistical power while reducing experimental cost (Kohavi et al. 2009; Barajas et al. 2016; Johnson et al. 2016b, 2017b,a). Design feature options include the randomization method (e.g., randomize before targeted users are selected or after, sample sizes for the control and treatment groups,²⁰ type of ads to be shown to the control group [e.g., placebo ad or competing advertiser’s ad],²¹ length of the experiment(s), and sample data period [i.e., data collection period prior and/or post experiment]).²²

Using the results from the experiments, advertisers ultimately want to infer the relative effectiveness of different ad buying strategies and optimally re-allocate resources. For example, Hoban Bucklin (2015) calculate marginal effects and elasticities along the purchase funnel, and suggest ways to re-allocate impressions across users at different stages. Similarly, Barajas et al. (2016) compare performance-based campaign with CPM-based campaign and leverage the results in re-selecting target users. Presumably advertisers seek to test different ad designs and inventory characteristics, targeting types, targeting algorithms, frequency and timing of ad exposures, and so forth. Based on the result from the previous experiment, how should the next experiment be designed? Related questions are in which order these decision variables need to be tested and which variables should be grouped together in the testing environment to reduce the experimental cost. There exists an

²⁰ Equal-sized A/B testings improve statistical power (Lewis et al. 2015), while dynamically assigning more users toward the better performing choices can potentially reduce experimental costs (Scott 2010; Schwartz et al. 2017). Optimally balancing this trade-off is an important area for future research.

²¹ Johnson et al. (2016b) show that adding “control ads” to the control group, which can record those users in the control group who would have received the treatment ad had they been in the treatment group, can improve statistical precision. Johnson et al. (2017b) develop “ghost ads” methodology that can reduce experimental cost while attaining the precision when running control ads.

²² Lewis et al. (2015) point out the trade-off in picking a cut-off date for the sample data. While a wider window of time can potentially capture the long-term effect, the amount of noise tends to increase faster than the signal (treatment effect) as one moves further from the campaign date.

increasing demand for automated A/B testing solutions, and more work would be welcome in the context of designing and effectively learning through repeated testings toward finding the global optimal ad buying decision. Repeated feedback cycles (i.e., testing \rightarrow analyzing results \rightarrow updating \rightarrow re-testing) become more crucial when the display ad environment changes rapidly over time and when competitors’ ad strategies change in response to focal advertiser’s re-optimization.

3.3.3 Advertising Effect and Privacy

More precise targeting generally increases ad effectiveness by delivering ads that are in consumers’ interests (Beales 2010; Goldfarb Tucker 2011b; Farahat Bailey 2012; Lambrecht Tucker 2013; Bleier Eisenbeiss 2015). However, finer targeting requires some degree of privacy intrusion, and may backlash as customers’ sense of vulnerability and privacy concerns increase (Goldfarb Tucker 2011a; Tucker 2014b; Aguirre et al. 2015).²³ More stringent privacy policies (opt-in $<$ opt-out $<$ tracking ban) negatively impact profitability of advertisers and publishers, and there is still an on-going debate over regulatory controls (Johnson 2013; Johnson et al. 2017c). How privacy policies should be set in consideration with consumers’ welfare (benefits from increased protection of privacy, minus loss from being exposed to less relevant ads) and changes in welfares of advertisers, publishers, and intermediaries is an interesting area for future research.²⁴ Another important area for future research concerns how much data to use in targeting and the value of different types of data bought from third parties, when considering the trade-off between finer targeting and privacy concerns. For example, it is unclear whether and when advertisers benefit from acquiring information on users’ social graphs and leveraging it in behavioral targeting.

4 Publishers

A publisher on the supply side of the display ad ecosystem faces the following decisions: (i) pricing advertising inventory (§4.1), (ii) deciding what information to share and with whom (§4.2), and (iii) scheduling the ad delivery (§4.3).

²³ Readers are referred to Tucker (2012a), Goldfarb (2014), and Acquisti et al. (2016) for reviews on economics of privacy and targeted advertising.

²⁴ On a related note, progress is being made from a computational perspective in developing behavioral targeting systems that are privacy friendly. For example, Toubiana et al. (2010) build a behavioral profiling and targeting system that does not leak user information outside the web browser.

4.1 Setting Prices

Advertisers jointly set prices conditioned on advertiser valuations across the guaranteed selling channel and the RTB channel. However, most prior literature explores publisher pricing decisions conditional on the selling channel, and we address some challenges inherent in the joint pricing problem.

4.1.1 Inferring Advertiser Valuations

It is self-evident that a first step in setting prices across channels is understanding advertiser valuations as discussed in sections 3.2.3 and 3.3. That said, publishers face additional challenges in ascertaining their respective advertisers' valuations as publishers do not observe advertisers' proprietary data (e.g., purchase conversion rate). Eliciting advertisers' valuations based on the limited information available to publishers, and how these may vary across selling channels, is an important area for future research. For example, joint consideration of the bid amount in the exchange auction and the fixed price paid in the guaranteed contracts and the volume bought in each selling channel can provide more information about advertisers' valuations.

4.1.2 Pricing in Guaranteed Selling Channel

In the guaranteed selling channel, publishers and advertisers come to an agreement on a fixed price pertaining to when, where, and how the ads will be displayed. We first discuss the most common CPM pricing, then explore other forms of pricing mechanisms.

CPM Pricing The difficulty in committing to guaranteed contracts in advance is that publishers face uncertain demand (advertisers' arrivals) and ad inventory (viewers' arrivals) (Feige et al. 2008). To address these uncertainties, existing research typically constructs a revenue management model in which advertisers arrive sequentially with desired quantities of impressions and valuations, and in which the publisher makes contractual decisions prior to knowing other demands arriving in the future. The publisher either (i) determines whether or not to accept arriving advertiser's proposed price and quantity, where the accepted request may be canceled at a cost (Babaioff et al. 2009; Constantin et al. 2009; Roels Fridgeirsdottir 2009), or (ii) dynamically determines a price to quote to an arriving advertiser's demand (Fridgeirsdottir Najafi-Asadolahi 2016).

The actual guaranteed contractual process more closely resembles a bargaining process, where the outcome may depend on several factors including advertiser's willingness to pay, publisher's

outside option, competition, business relationship, and negotiation skills. All these factors suggest interesting questions for research. As an example, little is known about the role sales people play in selling guaranteed contracts. Sales people might lower the search costs of advertisers in finding the right website, or convey additional information and value about the impressions to the advertisers. With the increase in programmatic direct, the role and value of sales people may change.

Finally, as advertisers have different valuations for varying inventory characteristics (e.g., size, format), publishers can bundle some of these ad units (e.g., desktop+mobile) to better price discriminate. Which ad units to bundle at which fixed bundle price is an open question.

Other Pricing Schemes Although CPM pricing is predominant in display ad, publishers can also consider other pricing schemes such as CPC (Mangani 2004; Fjell 2009; Asdemir et al. 2012; Najafi-Asadolahi Fridgeirsdottir 2014) or subscription fees (Baye Morgan 2000; Kumar Sethi 2009). Kumar Sethi (2009), for example, develop a dynamic hybrid revenue model that determines the optimal subscription fee level and amount of ad on a website over time, in consideration with the cost of serving each customer, cost of presenting ads, and cost of changing the subscription fee. Another interesting approach is taken in Goldstein et al. (2015), in which a time-based pricing scheme is considered because duration of an ad affects memory retention. Future research is warranted in exploring and comparing different pricing mechanisms that further look at the effect of advertisers’ competition, publishers’ competition, and consumers’ multi-homing behaviors across websites.

4.1.3 Pricing in Non-Guaranteed Selling Channel

Compared to the large body of literature focusing on advertisers’ optimal bidding strategies, research on publishers’ optimal pricing policies in RTB are relatively scarce. As in guaranteed contracts, information asymmetry (Hu et al. 2016) or double marginalization (Dellarocas 2012) can affect the publisher’s choice of pricing type in RTB, i.e., CPM, CPC, or other cost-per-action (CPA) auctions. Existing literature further considers publisher’s incentives in adopting a hybrid format where advertisers are allowed to choose between CPM or CPC type auctions (Zhu Wilbur 2011; Liu Viswanathan 2013).²⁵

²⁵ Facebook is one example where advertisers can choose to bid based on CPM, CPC, or other conversion metrics (e.g., Page Likes).
<https://adespresso.com/blog/everything-need-know-facebook-ads-bidding/>

While intermediaries typically control many aspects of the auction mechanism (e.g., second-price vs. first-price auction format, see § 5.3), the publisher still faces its pricing decision conditioned on the auction rules. For example, adding reserve prices in RTB auctions can substantially increase publishers’ revenue, especially when asymmetric, heterogeneous advertisers compete in thin auctions (Yuan et al. 2014; Paes Leme et al. 2016). The optimal reserve prices would need to take into account that advertisers’ buying strategies reflect some practical constraints they face such as budget (Balseiro et al. 2015) or reach (Choi Mela 2018). When the distribution of bids is not known by the publisher, Amin et al. (2013), Medina Mohri (2014), and Cesa-Bianchi et al. (2015) explore how to learn optimal reserve prices in a non-parametric fashion. An open issue is setting the optimal reserve prices in the face of competing publishers whose inventories are considered as (imperfect) substitutes by advertisers, and to develop algorithms to deploy the reserve prices at scale into ad exchange auctions.

Private marketplaces in the non-guaranteed selling channel are becoming increasingly popular and differ from open auctions inasmuch as only select advertisers can participate.²⁶ Private marketplaces typically consist of pre-negotiated, yet non-guaranteed, contracts for impressions. A common format is to provide “priority” advertisers access to invitation-only auctions inventory at a fixed CPM prior to impression being sold in the open auction.²⁷ Despite the added complexity of this format relative to open auctions, advertisers are asking for private deals to reduce brand safety concerns, to ensure a wider reach of impressions, and to access a coveted premium audience data (without having to commit to upfront contracts). Publishers often provide valuable audience segment targeting based on their own proprietary data for preferred deals, but not for open auctions. Different levels of information revealed can be a powerful product differentiator and publishers can improve revenue with these deals (Chen Wang 2015; Mirrokni Nazerzadeh 2017). However, a long list of questions remain regarding this nascent and growing pricing format, including when/how to offer these deals and at what prices, to whom to give private auction access, and how to set reserve prices in private auctions when the unsold impression can subsequently be sold at an open auction.

²⁶ <https://www.emarketer.com/newsroom/index.php/us-direct-and-private-marketplaces-take-increasing-share-of-programmatic/>

²⁷ <https://www.entrepreneur.com/article/322324>

4.1.4 Interrelation of Guaranteed and Non-Guaranteed Selling Channels

As advertisers can enter in both selling channels and publishers can sell in both, pricing should ideally be optimized jointly (Chen et al. 2014). When publishers are considering the inter-related benefits and costs of selling in both channels, it remains unclear what the jointly optimal guaranteed fixed prices and reservation prices should be. Moreover, other pricing mechanisms might yield even higher joint profits for publishers in this dual-selling channel structure by further price discriminating based on advertisers' valuations or uncertainty. For example, a publisher can consider selling call-options that can be transferred among advertisers, who may pay a premium to hedge against fluctuating prices in real-time auctions.

4.2 Information

What information to share with the advertisers and at what price are important questions for the publishers and warrant further studies. In the guaranteed selling channel, the publisher faces a trade-off in providing more information to the advertisers and making more granular level targeting options available. On the one hand, the publisher can facilitate price discrimination as advertisers are willing to pay a premium for an additional layer of targeting. On the other hand, the publisher can choose to disclose less information by bundling impressions together. When the publisher sells a large number of ad impressions and does not know advertisers' heterogeneous valuations for these impressions, predicting advertisers' valuations for a bundle of goods is easier because, from the law of large numbers, the distribution of valuations for the bundle tends to concentrate around its mean (see, e.g., Bakos Brynjolfsson 1999). Better predictions can, in turn, lead to higher revenues for the publisher. Additionally, coarser targeting criteria can reduce the risk of not meeting the contractual terms and paying a penalty in the event that the publisher falls short of the impressions.²⁸ Thus, the publisher accounts for the benefits and costs of providing more information to the advertiser when underwriting the targeting terms in guaranteed contracts.

Similarly in RTB, the publisher may benefit from providing more information as advertisers bid higher amounts conditional on participating in the auctions (De Corniere De Nijs 2016; Hummel McAfee 2016; Ada et al. 2018). At the same time, providing more information decreases the

²⁸ For example, some advertisers may value and want to target first-time home buyers. First-time home buyers in their 30s visiting The New York Times will be a smaller set than users in the 30s, and the visits of the former group on a given day will likely exhibit higher volatility.

number of participating bidders and creates thin markets, as fewer advertisers are interested in a given impression with highly differentiated attributes (Levin Milgrom 2010; Fu et al. 2012; Chen Stallaert 2014). In the presence of advertisers’ and consumers’ heterogeneous preference for targeting, Gal-Or et al. (2018) show that the differentiation in targeting levels offered to advertisers in equilibrium depends on the intensity of competition among platforms who face multi-homing consumers.

Providing more information, thus finer targeting ability, also results in displaying ads that highly coincide with users’ recent behaviors, which may increase privacy concerns and reduce ad responses (Lu Yang 2015). Furthermore, advertisers can use shared information to target the same user on other, cheaper websites (Ghosh et al. 2015). Because of these trade-offs, a publisher’s decision to share (Emek et al. 2014) or to price and sell his/her own proprietary data is more complex than that of a third party selling cookies (Bergemann Bonatti 2015). This raises questions regarding whether publishers should share information about users’ behaviors prior to and/or after the ad exposures, and whether it will benefit publishers to share KPIs for each specific type of ad inventory or as overall averages.

4.3 Delivery of Ad and Impressions

Ad scheduling is an inherently challenging task for the publisher. On one hand, publishers have guaranteed contracts to fulfill. On another hand, publishers can potentially enjoy high bids from those advertisers who want to cherry pick impressions in RTB. In addition, there is an uncertainty about the supply of impressions which are largely determined by users’ browsing behaviors.

For the guaranteed selling channel, prior research has considered allocating arriving impressions among contracted advertisers, seeking to maximize publisher revenue (Adler et al. 2002; Kumar et al. 2006; Feldman et al. 2009; Vee et al. 2010; Devanur et al. 2011; Ciocan Farias 2012; Turner 2012) and/or ad responses (Korula et al. 2013; Chickering Heckerman 2003) when contracts differ in prices, campaign schedules, targeting criteria, and delivery status. As RTB markets have grown, there has been increased interest in dynamic allocation, where publishers first request bids for each impression (often by adopting header bidding) and compare the winning price to the option value of assigning the impression to the best matching guaranteed contract (Ghosh et al. 2009a; Balseiro et al. 2014; Arnosti et al. 2016; Chen 2017; Sayedi 2018). These studies show that dynamic

allocation can yield higher profits for the publishers than first sending (random quality) impressions to the guaranteed contracts. Nevertheless, many publishers still prefer selling premium high quality inventories (e.g., front page, leaderboard) via guaranteed contracts, and limit the application of dynamic allocation to lower quality inventories with lower guaranteed contract prices. Presumably for certain types of inventories, leveraging advertisers’ willingness to pay high premium for reach or brand safety (ensured via guaranteed contracts) benefit publishers more than leveraging high bids for cherry picked impressions in RTB. In sum, allocation algorithms remain an active area for inquiry and are likely incumbent upon how advertisers value upfront contracts.

In addition, publishers could allow budget pacing (Balseiro et al. 2017b; Conitzer et al. 2017), impression pacing (Bhalgat et al. 2012), or frequency capping (Hojjat et al. 2017), and assess how these options provided to advertisers affect publisher’s optimal ad delivery plan. Another question relates to controlling supply of impressions to facilitate ad delivery or to raise advertiser competition. Publishers could consider buying additional traffic (e.g., via Facebook ad) to intentionally increase supply, or intentionally not sell certain ad spaces at a given point in time. Varying (or smoothing) supply of impressions affects advertiser competition and allocation of impressions among advertisers, making this another potential direction for future research.

5 Intermediaries

As outlined in the display ad market ecosystem discussion, many intermediaries (DSPs, SSPs, ad exchanges, ad networks, data aggregators) serve technological needs in RTB and facilitate the match between advertisers and publishers. These intermediaries affect market outcomes, yet research remains sparse regarding the implications of intermediaries’ decisions. We discuss (i) welfare implications of enabling technologies (§5.1), (ii) principal-agent incentive mis-alignments (§5.2), (iii) intermediaries’ decisions regarding pricing (§5.3), and (iv) whether competing intermediaries can benefit from sharing information such as bids and payments (§5.4).

5.1 Enabling Technologies and Vertical Integration

The evolution of buying and selling practices in display advertising are tightly integrated with the development of intermediaries’ enabling technologies (e.g., data collection, data analysis, RTB, real-time ad serving, A/B testing and optimization tools). With such technologies, intermediaries can better match publishers’ impressions (consumers) to more relevant advertisers (Broder et al.

2007; Chakrabarti et al. 2008; Zhang Katona 2012). Further, Balseiro Candogan (2017) show that intermediation leads to an increase in the overall market efficiency when optimal contracts are offered to budget-constrained advertisers with private information on budgets and targeting criteria.

On the other hand, these intermediaries charge high fees for their services. Although the fees charged by these ad tech intermediaries vary widely depending on the level of services provided in buying and selling ads, their revenues constitute approximately 55% of the ad spend in RTB.²⁹ High intermediation fees have led a large fraction of brand advertisers to execute programmatic buying in-house.³⁰ However, high costs associated with sourcing highly-skilled media procurement talent to manage the time and effort required in buying ads has led some advertisers to turn back to their agencies for help.³¹ Overall, the welfare implications of intermediaries who enable new practices in display ad markets (e.g., existence of RTB and targeted ad buying at scale) should be examined.³²

5.2 Agency

The existence of intermediaries (with their enabling technologies) raises interesting research questions as economic incentives of principals and intermediaries may not be aligned. Balseiro et al. (2017a) explore how the structure of intermediation network affects the profits of participants when advertisers' values are private, and show that intermediaries have incentives to shade bids and not to allocate impressions, even when profitable for their downstream customers. Allouah Besbes (2017), on the other hand, explore the welfare implications of the collusive behavior of DSPs who represent multiple advertisers, and show that moving from collusion (DSP submitting a single bid to limit competition among the advertisers it represents on a given impression) to non-collusion (DSP bidding for each advertiser independently of other advertisers it represents) leads to a Pareto improvement when taking into account the publisher's reaction such as increasing reserve prices. Future research is called for looking at different objectives of intermediaries, that might differ

²⁹ <https://www.iab.com/insights/iab-programmatic-revenue-report-2014-results/>

³⁰ According to the IAB, 65% of brand advertisers surveyed have fully or partially moved programmatic buying functions in-house. https://www.iab.com/wp-content/uploads/2018/05/IAB-Programmatic-In-Housing-Whitepaper_v5.pdf

³¹ <https://digiday.com/marketing/vodafone-backtracks-programmatic-ad-buying/>

³² Although in a different industry, Salz (2017) explains how the existence of intermediaries improves overall welfare by reducing incurred search costs and reallocating contracts to lower cost carters in New York City trade-waste market.

inherently from those of the principals, and appropriate contract mechanism design.

5.3 Pricing Mechanisms

In this sub-section, we first discuss existing considerations regarding intermediaries’ pricing (i.e., auction mechanism design) on revenues of advertisers and publishers. We then outline new developments and pricing challenges (such as those associated with dynamic auction design and header bidding).

One concern spanning the existing literature is that ad impressions are highly differentiated, valuations are heterogeneous across advertisers, and the distribution of advertisers’ valuations often violates the regularity conditions in Myerson (1981).³³ Moreover with user-specific information available in RTB, it is likely that advertisers have private information suggesting that certain types of impressions are commonly valued (Abraham et al. 2016). Ad exchanges further face asymmetric information states regarding publishers’ opportunity cost of selling the impressions and competition (Gomes Mirrokni 2014). With such characteristics, the Revenue Equivalence Theorem fails to hold, and ad exchanges can improve their revenues by designing mechanisms that are different from the commonly used second-price, sealed-bid auction. For example, Celis et al. (2014) proposes a randomized sales mechanism that outperforms an optimal second-price auction, where the buyer can choose to “buy-it-now” at a fixed price or “take-a-chance” in an auction with random allocation among top $d > 1$ bidders.

Going forward, there has been an increase in interest regarding the design of dynamic auctions. Because most advertisers participate in multiple auctions throughout their campaign, the auctioneer can offer dynamic mechanisms that link incentives across auctions and improve revenues compared to repeated static auctions (Balseiro et al. 2017c; Mirrokni et al. 2018). Further, the commission structure among intermediaries is likely to affect the intermediaries’ incentives, the auction formats they adopt, and the associated market outcomes. More attention is warranted to ascertain this dependency.

Competition among advertisers and publishers with multiple (competing) intermediaries’ pose additional complexity in display markets. One recently introduced coordination mechanism is

³³ In analyzing the second-price auction data from Microsoft Advertising Exchange, Celis et al. (2014) find that there is often a large gap between the highest and second-highest bid, and that the gap is bigger than the winning payment price on average.

header bidding, which gathers bids from multiple ad exchanges.³⁴ Although different sets of advertisers (DSPs) interact with different ad exchanges, header bidding enables publishers to coalesce demand for the inventory and increase revenue by intensifying competition among ad exchanges and advertisers.³⁵

Header bidding can be construed as a two-stage sequential game: in the second stage ad exchanges independently run second-price auctions, and in the first stage the publisher chooses to deliver the impression to the ad exchange paying the highest. The outcome from this game can drastically differ from that from a single second-price auction with all advertisers' bids on one intermediary platform. For example, the highest bidder is not guaranteed to win the impression because most intermediaries submit their second-highest bids (=payments under second-price auctions).³⁶ To ensure that the highest bidder wins the impression and to increase transparency in payments, many ad exchanges are now instead adopting first-price auctions (or soft floors, Zeithammer 2016), but it is not clear whether advertisers/publishers will benefit or hurt by this shift.^{37 38} In spite of the rapid growth in header bidding, research is virtually non-existent regarding how advertisers compete in multi-stage auction game with intermediaries competing toward increasing their own commissions (on payments by advertisers), and the attendant impact on the welfares of advertisers, publishers, and intermediaries.

5.4 Information

Section 4.2 discusses the publisher's problem with regard to information, and an analogous problem exists for intermediaries. Intermediaries possess information spanning advertisers and/or publishers and can therefore assess the competitive landscape more completely than a single advertiser or publisher. DSPs have access to ad budgets, bidding strategies, winning bids, and payments across advertisers. SSPs and ad networks have information regarding how impressions are served among advertisers/campaigns and associated ad responses across publishers. Ad exchanges have information about all winning and losing bids from advertisers and inventory characteristics from

³⁴ <https://www.sovrn.com/blog/header-bidding-grows-up/>

³⁵ Prior to the adoption of header bidding, publishers would avail an incoming impression at the primary ad exchange, and if not sold (e.g., the highest bid falls below the reserve price), then this impression would subsequently be auctioned at the next ad exchange and so on (i.e., moving down the waterfall sequentially until sold).

³⁶ <https://www.hearts-science.com/sold-for-more-than-you-should-have-paid/>

³⁷ <https://adage.com/article/digital/google-adx-moving-a-price-auction/316894/>

³⁸ According to eMarketer, the use of first (second) price auction format increased (decreased) from 5.8% (75.1%) in 12/2017 to 43.4% (32.8%) in 03/2018.

<https://www.emarketer.com/content/first-price-auctions-are-driving-up-ad-prices>

publishers.

Limited data available to a single advertiser or publisher raises questions of how much/when/under which circumstances information should be shared between intermediaries and advertisers (and also publishers). For example, Sun et al. (2016) consider the timing of ad exchanges' revelation to advertisers and find that hiding impression-level information at bid-time, but revealing at win-time, can increase the seller's revenue, especially when the gap between the highest and second-highest bid is large. In addition, ad exchanges can consider providing bid landscape reports to advertisers (and/or publishers) to help them better predict the distribution of highest bids. Disclosing bid information to advertisers (or publishers) can range from revealing all bids and identity of the bidders to revealing simple summary statistics as quantiles. Advertisers may or may not consent to such disclosure and receiving more information about their competitors.

Intermediaries may have different data-sharing incentives from advertisers and publishers (Marotta et al. 2017). Rafieian Yoganarasimhan (2018b) consider the data-sharing arrangements between the ad network and advertisers, and find that ad networks have the incentive to withhold information to increase advertiser competition and increase ad networks' revenue. Resolving the information asymmetry between advertisers and publishers by the intermediary will further depend on the competition among intermediaries. For example, header bidding will change the incentives of intermediaries in data sharing arrangements as competition among ad exchanges intensifies.

Another avenue for research is the market for data. Advertisers have different types of proprietary data in house and often consider buying third-party data for additional user information. For example, Amazon has access to users' browsing and purchase histories on its site, whereas P&G might have a harder time linking online users to their offline purchases. These different levels of user information raises the question of when it is beneficial for the advertisers to buy additional data and how the intermediary should price the information (Bhawalkar et al. 2014; Dalessandro et al. 2014). Related questions of interest pertain to the types of data third-parties should collect that can best augment advertisers' own proprietary data toward improving consumers' ad experiences and privacy concerns. The information intermediaries share with advertisers or publishers can affect bargaining power among these players, so it would be desirable to address this problem as well.

6 Display Advertising Contexts

In addition to the more general display advertising considerations we note above, various contextual factors are becoming more germane as technology in this space evolves. Specifically, we address mobile and social media in this section, and discuss video ads, transparency related issues, a brief history of display ads, and sponsored search advertising in the Online Appendix.

6.1 Mobile

Various aspects of the display markets in mobile advertising warrant additional consideration owing to its rapid growth and unique ecosystem.³⁹ First, search costs are higher (Ghose et al. 2012) and the information-carrying capacity is more constrained on mobile as screen sizes are smaller. As such, Bart et al. (2014) show that mobile display advertising is more effective for higher involvement and utilitarian (vs. lower involvement or hedonic) products.

A second key point of difference is that mobile offers a wide new array of targeting variables that can be used for the purchase of impressions. Ghose (2017) enumerates context (e.g., Ghose et al. 2019a), location (e.g., Fang et al. 2015; Fong et al. 2015; Dubé et al. 2017), time (e.g., Danaher et al. 2015), saliency (e.g., Ghose et al. 2012), crowdedness (e.g., Andrews et al. 2015), weather (e.g., Li et al. 2017), historical shopping patterns/trajectory (e.g., Ghose et al. 2019b), social dynamics (e.g., Provost et al. 2015; Ghose et al. 2017), and tech mix as additional factors that moderate the efficacy of advertising. The value of this information for targeting is exemplified by Andrews et al. (2015), who demonstrate that users are responsive to marketing offers in crowded environments, and Luo et al. (2013) who detail the interaction between proximity to the focal venue and the time a promotion is received. An active area of research is how this rich information can be exploited in the mobile display ad ecosystem and the resulting implication for the players in this market.

A third unique aspect pertains to balancing advertising across mobile web and in-app.⁴⁰ Although mobile apps consume nearly 90% of the time adults’ spend on their smartphone (including non-ad supportive apps such email apps), there exists a greater diversity of publishers and audiences using mobile web. Accordingly, coordinating campaigns across apps and browsers requires additional consideration, though there is little research to date on understanding consumers’ inter-

³⁹ See Grewal et al. (2016) and Ghose (2017) for discussions on mobile advertising, including display ads.

⁴⁰ <https://www.iab.com/events/iab-make-mobile-work-webinar-in-app-mobile-web-confusion/>

actions (e.g., in-app browsing) and advertising effectiveness across these channels.⁴¹

Fourth for the in-app advertising, there exist multiple platforms (e.g., iOS, Android). As users and apps differ significantly across platforms, architecting advertiser campaigns across multiple platforms is an open issue. On the supply- and intermediary-sides, open questions include the role of app providers (who are often both publishers and advertisers, Lee et al. 2017b), tying information across platforms or ad networks (for example, to enable in-app parallel bidding),⁴² and optimal ad-sequencing (e.g., in-app fading ads in Sun et al. 2017; Rafeian Yoganarasimhan 2018a).

Finally, the concept of mobile advertising engagement is expanding and evolving. For example, information regarding consumers’ emotions from facial expressions (e.g., face unlock) or biometrics (e.g., watch) will likely affect advertising strategies. Likewise, concerns remain regarding how advertisers should think about audio advertising as consumers are increasingly turning to smart speakers for convenient access to news and entertainment.

6.2 Social Media

As with mobile advertising, social media advertising enables advertisers to tap into a richer set of variables for targeting users, such as the network structure itself and its users’ potential sharing and liking actions. Moreover, because almost all users are registered, social sites have a large set of user descriptors (Gordon et al. 2016). Finally, targeting in this context relies not only on the response to a single user, but the ripple effect of social influence (Zubcsek Sarvary 2011; Bakshy et al. 2012; Tucker 2012b, 2014a; Lee et al. 2017a; Zhang et al. 2016a). Little research exists on bidding in ad markets where user valuations are interdependent.

Perhaps more germane is social advertising, which is dominated by Facebook. It is largely vertically integrated, so its ad buying network is vastly different than that characterized in Section 2. Not only does Facebook serve as a SSP, DMP, and Ad Exchange, but it is also a publisher. As a result, Facebook can more easily incorporate its own data in designing the auction mechanism and has greater control over balancing the needs of its viewers and the needs of its advertisers. For example, Facebook might accept lower bids in order to obtain higher reader engagement since showing ads that are not well aligned with reader interests can induce its viewers to leave the site.⁴³

⁴¹ <https://www.emarketer.com/Article/Comparing-Responses-In-App-Mobile-Web-Ads/1016037>

⁴² <https://pubmatics.com/blog/in-app-header-bidding/>

⁴³ Facebook currently adopts a *modified* version of second-price auction in which two factors, other than bid, are also considered in determining the winner: (i) estimated action rates and (ii) ad quality and relevance factors (<https://www.facebook.com/business/help/430291176997542>). This modified auction format may deliver more rele-

Trading off the relative weight of advertiser and user interests in determining the auction winner would be a beneficial avenue for research in both traditional (Mookerjee et al. 2016; Stourm Bax 2017) and social media display markets.⁴⁴

7 Conclusion

The display ad market is economically substantial and rapidly growing. The growth is largely driven by RTB, in which many intermediaries exist to serve the technological needs of advertisers and publishers. The guaranteed selling channel and RTB have distinguishing characteristics, especially in terms of players involved, information available, and pricing mechanisms. Despite some industry experts' belief that all ads will be bought and sold via RTB, the dual selling channels will likely co-exist in the long-run as advertisers and publishers have incentives to leverage the distinctive features of both channels (Athey et al. 2017; Sayedi 2018).

In this article, the ecosystem of the display ad market, with dual selling channels, is first outlined. Then, for each player involved (advertisers, publishers, and intermediaries), papers are reviewed across disciplines pertaining to each player's decisions. Many open inter-disciplinary research questions remain to be explored, especially in joint consideration of guaranteed and non-guaranteed selling channels, and we propose potential directions to stimulate fruitful research in this field.

vant ads to consumers cost-effectively.

⁴⁴ Of note, vertical integration considerations are extending to other contexts. AT&T's recent purchase of Time Warner Inc. is an attempt to bundle ad network (distribution) and content.

References

- Abhishek V, Fader P, Hosanagar K (2012) Media exposure through the funnel: A model of multi-stage attribution Working paper.
- Abraham I, Athey S, Babaioff M, Grubb MD (2016) Peaches, lemons, and cookies: Designing auction markets with dispersed information Working paper.
- Abraham M (2008) The off-line impact of online ads. *Harvard Business Review* 86(4):28.
- Acquisti A, Taylor C, Wagman L (2016) The economics of privacy. *Journal of Economic Literature* 54(2):442–492.
- Ada S, Abou Nabout N, Feit EM (2018) Disclosing information about ad placements in online display advertising auctions Working paper.
- Adler M, Gibbons PB, Matias Y (2002) Scheduling space-sharing for internet advertising. *Journal of Scheduling* 5(2):103–119.
- Agarwal A, Chapelle O, Dudík M, Langford J (2014) A reliable effective terascale linear learning system. *Journal of Machine Learning Research* 15(1):1111–1133.
- Agarwal D, Agrawal R, Khanna R, Kota N (2010) Estimating rates of rare events with multiple hierarchies through scalable log-linear models. *Proceedings of the 16th international conference on knowledge discovery and data mining*, 213–222 (ACM).
- Aguirre E, Mahr D, Grewal D, de Ruyter K, Wetzels M (2015) Unraveling the personalization paradox: The effect of information collection and trust-building strategies on online advertisement effectiveness. *Journal of Retailing* 91(1):34–49.
- Ahmed A, Low Y, Aly M, Josifovski V, Smola AJ (2011) Scalable distributed inference of dynamic user interests for behavioral targeting. *Proceedings of the 17th international conference on knowledge discovery and data mining*, 114–122 (ACM).
- Aksakalli V (2012) Optimizing direct response in internet display advertising. *Electronic Commerce Research and Applications* 11(3):229–240.
- Allouah A, Besbes O (2017) Auctions in the online display advertising chain: A case for independent campaign management Working paper.
- Aly M, Hatch A, Josifovski V, Narayanan VK (2012) Web-scale user modeling for targeting. *Proceedings of the 21st international conference on world wide web*, 3–12 (ACM).
- Amin K, Rostamizadeh A, Syed U (2013) Learning prices for repeated auctions with strategic buyers. *Advances in Neural Information Processing Systems*, 1169–1177.
- Andrews M, Luo X, Fang Z, Ghose A (2015) Mobile ad effectiveness: Hyper-contextual targeting with crowdedness. *Marketing Science* 35(2):218–233.
- Arnosti N, Beck M, Milgrom P (2016) Adverse selection and auction design for internet display advertising. *The American Economic Review* 106(10):2852–2866.
- Asdemir K, Kumar N, Jacob VS (2012) Pricing models for online advertising: Cpm vs. cpc. *Information Systems Research* 23(3-part-1):804–822.
- Athey S, Calvano E, Gans JS (2014) The impact of the internet on advertising markets for news media Working paper.
- Athey S, Calvano E, Gans JS (2017) The impact of consumer multi-homing on advertising markets and media competition. *Management Science*, forthcoming .
- Babaioff M, Hartline JD, Kleinberg RD (2009) Selling ad campaigns: Online algorithms with cancellations. *Proceedings of the 10th conference on electronic commerce*, 61–70 (ACM).
- Bagherjeiran A, Parekh R (2008) Combining behavioral and social network data for online advertising. *International conference on data mining workshops*, 837–846 (IEEE).
- Bakos Y, Brynjolfsson E (1999) Bundling information goods: Pricing, profits, and efficiency. *Management Science* 45(12):1613–1630.
- Bakshy E, Eckles D, Yan R, Rosenn I (2012) Social influence in social advertising: Evidence from field experiments. *Proceedings of the 13th conference on electronic eommerce*, 146–161 (ACM).
- Balseiro S, Candogan O, Gurkan H (2017a) Multi-stage intermediation in display advertising Working paper.
- Balseiro S, Gur Y (2017) Learning in repeated auctions with budgets: Regret minimization and equilibrium Working paper.
- Balseiro S, Kim A, Mahdian M, Mirrokni V (2017b) Budget management strategies in repeated auctions. *Proceedings of the 26th International Conference on World Wide Web*, 15–23 (International World Wide Web Conferences Steering Committee).

- Balseiro SR, Besbes O, Weintraub GY (2015) Repeated auctions with budgets in ad exchanges: Approximations and design. *Management Science* 61(4):864–884.
- Balseiro SR, Candogan O (2017) Optimal contracts for intermediaries in online advertising. *Operations Research* 65(4):878–896.
- Balseiro SR, Feldman J, Mirrokni V, Muthukrishnan S (2014) Yield optimization of display advertising with ad exchange. *Management Science* 60(12):2886–2907.
- Balseiro SR, Mirrokni VS, Leme RP (2017c) Dynamic mechanisms with martingale utilities. *Management Science* 64(11):5062–5082.
- Barajas J, Akella R, Holtan M, Flores A (2016) Experimental designs and estimation for online display advertising attribution in marketplaces. *Marketing Science* 35(3):465–483.
- Bart Y, Stephen AT, Sarvary M (2014) Which products are best suited to mobile advertising? a field study of mobile display advertising effects on consumer attitudes and intentions. *Journal of Marketing Research* 51(3):270–285.
- Baye MR, Morgan J (2000) A simple model of advertising and subscription fees. *Economics Letters* 69(3):345–351.
- Beales H (2010) The value of behavioral targeting. *Network Advertising Initiative* .
- Bergemann D, Bonatti A (2015) Selling cookies. *American Economic Journal: Microeconomics* 7(3):259–294.
- Berman R (2015) Beyond the last touch: Attribution in online advertising Working paper.
- Bhalgat A, Feldman J, Mirrokni V (2012) Online allocation of display ads with smooth delivery. *Proceedings of the 18th international conference on knowledge discovery and data mining*, 1213–1221 (ACM).
- Bhawalkar K, Hummel P, Vassilvitskii S (2014) Value of targeting. *International symposium on algorithmic game theory*, 194–205 (Springer).
- Bleier A, Eisenbeiss M (2015) Personalized online advertising effectiveness: The interplay of what, when, and where. *Marketing Science* 34(5):669–688.
- Braun M, Moe WW (2013) Online display advertising: Modeling the effects of multiple creatives and individual impression histories. *Marketing science* 32(5):753–767.
- Broder A, Fontoura M, Josifovski V, Riedel L (2007) A semantic approach to contextual advertising. *Proceedings of the 30th international conference on research and development in information retrieval*, 559–566 (ACM).
- Bruce NI, Murthi B, Rao RC (2017) A dynamic model for digital advertising: The effects of creative format, message content, and targeting on engagement. *Journal of Marketing Research* 54(2):202–218.
- Bucklin RE, Hoban PR (2017) Marketing models for internet advertising. *Handbook of Marketing Decision Models*, 431–462 (Springer).
- Cai H, Ren K, Zhang W, Malialis K, Wang J, Yu Y, Guo D (2017) Real-time bidding by reinforcement learning in display advertising. *Proceedings of the 10th international conference on web search and data mining*, 661–670 (ACM).
- Catalini C, Gans JS (2016) Some simple economics of the blockchain. Technical report, National Bureau of Economic Research.
- Celis LE, Lewis G, Mobius M, Nazerzadeh H (2014) Buy-it-now or take-a-chance: Price discrimination through randomized auctions. *Management Science* 60(12):2927–2948.
- Cesa-Bianchi N, Gentile C, Mansour Y (2015) Regret minimization for reserve prices in second-price auctions. *IEEE Transactions on Information Theory* 61(1):549–564.
- Chakrabarti D, Agarwal D, Josifovski V (2008) Contextual advertising by combining relevance with click feedback. *Proceedings of the 17th international conference on world wide web*, 417–426 (ACM).
- Chapelle O (2015) Offline evaluation of response prediction in online advertising auctions. *Proceedings of the 24th international conference on world wide web*, 919–922 (ACM).
- Chapelle O, Manavoglu E, Rosales R (2015) Simple and scalable response prediction for display advertising. *Transactions on intelligent systems and technology* 5(4):61.
- Chatterjee P, Hoffman DL, Novak TP (2003) Modeling the clickstream: Implications for web-based advertising efforts. *Marketing Science* 22(4):520–541.
- Chen B, Wang J (2015) A lattice framework for pricing display advertisement options with the stochastic volatility underlying model. *Electronic Commerce Research and Applications*

- 14(6):465–479.
- Chen B, Yuan S, Wang J (2014) A dynamic pricing model for unifying programmatic guarantee and real-time bidding in display advertising. *Proceedings of the 8th international workshop on data mining for online advertising*, 1–9 (ACM).
- Chen J, Stallaert J (2014) An economic analysis of online advertising using behavioral targeting. *Mis Quarterly* 38(2):429–449.
- Chen M, Jacob VS, Radhakrishnan S, Ryu YU (2015) Can payment-per-click induce improvements in click fraud identification technologies? *Information Systems Research* 26(4):754–772.
- Chen Y, Pavlov D, Canny JF (2009) Large-scale behavioral targeting. *Proceedings of the 15th international conference on knowledge discovery and data mining*, 209–218 (ACM).
- Chen YJ (2017) Optimal dynamic auctions for display advertising. *Operations Research* 65(4):897–913.
- Chickering DM, Heckerman D (2003) Targeted advertising on the web with inventory management. *Interfaces* 33(5):71–77.
- Chiong K, Shum M, Webb R, Chen R (2018) Split-second decision-making in the field: Response times in mobile advertising. *Available at SSRN* .
- Choi H, Mela CF (2018) Display advertising pricing in exchange markets Working paper.
- Chung CY, Koran JM, Lin LJ, Yin H (2010) Model for generating user profiles in a behavioral targeting system US Patent 7,809,740.
- Ciocan DF, Farias V (2012) Model predictive control for dynamic resource allocation. *Mathematics of Operations Research* 37(3):501–525.
- Conitzer V, Kroer C, Sodomka E, Stier-Moses NE (2017) Multiplicative pacing equilibria in auction markets. *arXiv preprint arXiv:1706.07151* .
- Constantin F, Feldman J, Muthukrishnan S, Pál M (2009) An online mechanism for ad slot reservations with cancellations. *Proceedings of the 20th symposium on discrete algorithms*, 1265–1274 (Society for Industrial and Applied Mathematics).
- Crussell J, Stevens R, Chen H (2014) Madfraud: Investigating ad fraud in android applications. *Proceedings of the 12th international conference on mobile systems, applications, and services*, 123–134 (ACM).
- Cui Y, Zhang R, Li W, Mao J (2011) Bid landscape forecasting in online ad exchange marketplace. *Proceedings of the 17th international conference on knowledge discovery and data mining*, 265–273 (ACM).
- Dalessandro B, Hook R, Perlich C, Provost F (2015) Evaluating and optimizing online advertising: Forget the click, but there are good proxies. *Big Data* 3(2):90–102.
- Dalessandro B, Perlich C, Raeder T (2014) Bigger is better, but at what cost? estimating the economic value of incremental data assets. *Big Data* 2(2):87–96.
- Dalessandro B, Perlich C, Stitelman O, Provost F (2012) Causally motivated attribution for online advertising. *Proceedings of the 6th international workshop on data mining for online advertising and internet economy*, 7 (ACM).
- Danaher PJ (2007) Modeling page views across multiple websites with an application to internet reach and frequency prediction. *Marketing Science* 26(3):422–437.
- Danaher PJ, Lee J, Kerbach L (2010) Optimal internet media selection. *Marketing Science* 29(2):336–347.
- Danaher PJ, Smith MS, Ranasinghe K, Danaher TS (2015) Where, when, and how long: Factors that influence the redemption of mobile phone coupons. *Journal of Marketing Research* 52(5):710–725.
- Daswani N, Mysen C, Rao V, Weis S, Gharachorloo K, Ghosemajumder S (2008) Online advertising fraud. *Crimeware: understanding new attacks and defenses* 40(2):1–28.
- Datta A, Tschantz MC, Datta A (2015) Automated experiments on ad privacy settings. *Proceedings on privacy enhancing technologies* 2015(1):92–112.
- De Corniere A, De Nijs R (2016) Online advertising and privacy. *The RAND Journal of Economics* 47(1):48–72.
- Dellarocas C (2012) Double marginalization in performance-based advertising: Implications and solutions. *Management Science* 58(6):1178–1195.
- Devanur NR, Jain K, Sivan B, Wilkens CA (2011) Near optimal online algorithms and fast approximation algorithms for resource allocation problems. *Proceedings of the 12th conference on*

- electronic commerce*, 29–38 (ACM).
- Dubé JP, Fang Z, Fong N, Luo X (2017) Competitive price targeting with smartphone coupons. *Marketing Science* 36(6):944–975.
- Emek Y, Feldman M, Gamzu I, PaesLeme R, Tennenholtz M (2014) Signaling schemes for revenue maximization. *Transactions on Economics and Computation* 2(2):5.
- Evans DS (2008) The economics of the online advertising industry. *Review of Network Economics* 7(3).
- Evans DS (2009) The online advertising industry: Economics, evolution, and privacy. *The Journal of Economic Perspectives* 23(3):37–60.
- Fang Z, Gu B, Luo X, Xu Y (2015) Contemporaneous and delayed sales impact of location-based mobile promotions. *Information Systems Research* 26(3):552–564.
- Farahat A, Bailey MC (2012) How effective is targeted advertising? *Proceedings of the 21st international conference on world wide web*, 111–120 (ACM).
- Feige U, Immorlica N, Mirrokni V, Nazerzadeh H (2008) A combinatorial allocation mechanism with penalties for banner advertising. *Proceedings of the 17th international conference on world wide web*, 169–178 (ACM).
- Feily M, Shahrestani A, Ramadass S (2009) A survey of botnet and botnet detection. *3rd international conference on emerging security information, systems and technologies*, 268–273 (IEEE).
- Feldman J, Korula N, Mirrokni V, Muthukrishnan S, Pál M (2009) Online ad assignment with free disposal. *International workshop on internet and network economics*, 374–385 (Springer).
- Fjell K (2009) Online advertising: Pay-per-view versus pay-per-click - a comment. *Journal of Revenue and Pricing Management* 8(2-3):200–206.
- Fong NM, Fang Z, Luo X (2015) Geo-conquesting: Competitive locational targeting of mobile promotions. *Journal of Marketing Research* 52(5):726–735.
- Fridgeirsdottir K, Najafi-Asadolahi S (2016) Cost-per-impression pricing for display advertising Working paper.
- Fu H, Jordan P, Mahdian M, Nadav U, Talgam-Cohen I, Vassilvitskii S (2012) Ad auctions with data. *Conference on computer communications workshops*, 184–189 (IEEE).
- Gal-Or E, Gal-Or R, Penmetsa N (2018) The role of user privacy concerns in shaping competition among platforms. *Information Systems Research* 29(3):698–722.
- Ghose A (2017) *TAP: Unlocking the mobile economy* (MIT Press).
- Ghose A, Goldfarb A, Han SP (2012) How is the mobile internet different? search costs and local activities. *Information Systems Research* 24(3):613–631.
- Ghose A, Han S, Park S (2013) Analyzing the interdependence between web and mobile advertising: A randomized field experiment Working paper.
- Ghose A, Kwon HE, Lee D, Oh W (2019a) Seizing the commuting moment: Contextual targeting based on mobile transportation apps. *Information Systems Research* .
- Ghose A, Li B, Liu S (2017) Nudging mobile customers with real-time social dynamics. *Working paper* .
- Ghose A, Li B, Liu S (2019b) Mobile targeting using customer trajectory patterns. *Management Science* .
- Ghose A, Todri-Adamopoulos V (2016) Towards a digital attribution model: Measuring the impact of display advertising on online consumer behavior. *MIS Quarterly* 40(4):889–910.
- Ghosh A, Mahdian M, McAfee RP, Vassilvitskii S (2015) To match or not to match: Economics of cookie matching in online advertising. *Transactions on Economics and Computation* 3(2):12.
- Ghosh A, McAfee P, Papineni K, Vassilvitskii S (2009a) Bidding for representative allocations for display advertising. *International workshop on internet and network economics*, 208–219 (Springer).
- Ghosh A, Rubinstein B, Vassilvitskii S, Zinkevich M (2009b) Adaptive bidding for display advertising. *Proceedings of the 18th international conference on world wide web*, 251–260 (ACM).
- Goel S, Goldstein DG (2013) Predicting individual behavior with social networks. *Marketing Science* 33(1):82–93.
- Goldfarb A (2014) What is different about online advertising? *Review of Industrial Organization* 44(2):115–129.
- Goldfarb A, Tucker C (2011a) Online display advertising: Targeting and obtrusiveness. *Marketing Science* 30(3):389–404.

- Goldfarb A, Tucker C (2011b) Privacy regulation and online advertising. *Management science* 57(1):57–71.
- Goldfarb A, Tucker C (2014) Standardization and the effectiveness of online advertising. *Management Science* 61(11):2707–2719.
- Goldstein DG, McAfee RP, Suri S (2015) Improving the effectiveness of time-based display advertising. *Transactions on Economics and Computation* 3(2):7.
- Gomes R, Mirrokni V (2014) Optimal revenue-sharing double auctions with applications to ad exchanges. *Proceedings of the 23rd international conference on World wide web*, 19–28 (ACM).
- Gordon BR, Zettelmeyer F, Bhargava N, Chapsky D (2016) A comparison of approaches to advertising measurement: Evidence from big field experiments at facebook White paper.
- Grewal D, Bart Y, Spann M, Zubcsek PP (2016) Mobile advertising: A framework and research agenda. *Journal of Interactive Marketing* 34:3–14.
- Ha L (2008) Online advertising research in advertising journals: A review. *Journal of Current Issues & Research in Advertising* 30(1):31–48.
- He X, Pan J, Jin O, Xu T, Liu B, Xu T, Shi Y, Atallah A, Herbrich R, Bowers S, et al. (2014) Practical lessons from predicting clicks on ads at facebook. *Proceedings of the 8th international workshop on data mining for online advertising*, 1–9 (ACM).
- Hoban P, Arora N (2018) Measuring display advertising response using observational data: The impact of selection biases. *Available at SSRN 3264871*.
- Hoban PR, Bucklin RE (2015) Effects of internet display advertising in the purchase funnel: Model-based insights from a randomized field experiment. *Journal of Marketing Research* 52(3):375–393.
- Hojjat A, Turner J, Cetintas S, Yang J (2017) A unified framework for the scheduling of guaranteed targeted display advertising under reach and frequency requirements. *Operations Research* 65(2):289–313.
- Hu Y, Shin J, Tang Z (2016) Incentive problems in performance-based online advertising pricing: Cost per click vs. cost per action. *Management Science* 62(7):2022–2038.
- Hummel P, McAfee RP (2015) Loss functions for predicted click-through rates in auctions for online advertising. *Journal of Applied Econometrics*.
- Hummel P, McAfee RP (2016) When does improved targeting increase revenue? *Transactions on Economics and Computation* 5(1):4.
- Iyer K, Johari R, Sundararajan M (2014) Mean field equilibria of dynamic auctions with learning. *Management Science* 60(12):2949–2970.
- James GM, Paulson C, Rusmevichientong P (2018) Penalized and constrained optimization: An application to high-dimensional website advertising. Technical report, Discussion paper.
- Johnson G, Lewis RA, Nubbemeyer E (2017a) The online display ad effectiveness funnel & carry-over: Lessons from 432 field experiments Working paper.
- Johnson GA (2013) The impact of privacy policy on the auction market for online display advertising Working paper.
- Johnson GA, Lewis RA (2015) Cost per incremental action: Efficient pricing of advertising Working paper.
- Johnson GA, Lewis RA, Nubbemeyer EI (2017b) Ghost ads: A revolution in measuring ad effectiveness. *Journal of Marketing Research*.
- Johnson GA, Lewis RA, Reiley D (2016a) Location, location, location: Repetition and proximity increase advertising effectiveness Working paper.
- Johnson GA, Lewis RA, Reiley DH (2016b) When less is more: Data and power in advertising experiments. *Marketing Science* 36(1):43–53.
- Johnson GA, Shriver S, Du S (2017c) Consumer privacy choice in online advertising: Who opts out and at what cost to industry? Working paper.
- Jordan P, Mahdian M, Vassilvitskii S, Vee E (2011) The multiple attribution problem in pay-per-conversion advertising. *International symposium on algorithmic game theory*, 31–43 (Springer).
- Juan Y, Zhuang Y, Chin WS, Lin CJ (2016) Field-aware factorization machines for ctr prediction. *Proceedings of the 10th conference on recommender systems*, 43–50 (ACM).
- Kireyev P, Pauwels K, Gupta S (2016) Do display ads influence search? attribution and dynamics in online advertising. *International Journal of Research in Marketing* 33(3):475–490.
- Kohavi R, Longbotham R, Sommerfield D, Henne RM (2009) Controlled experiments on the web: Survey and practical guide. *Data mining and knowledge discovery* 18(1):140–181.

- Korula N, Mirrokni V, Nazerzadeh H (2016) Optimizing display advertising markets: Challenges and directions. *Internet Computing* 20(1):28–35.
- Korula N, Mirrokni VS, Zadimoghaddam M (2013) Bicriteria online matching: Maximizing weight and cardinality. *International conference on web and internet economics*, 305–318 (Springer).
- Kumar S, Jacob VS, Sriskandarajah C (2006) Scheduling advertisements on a web page to maximize revenue. *European journal of operational research* 173(3):1067–1089.
- Kumar S, Sethi SP (2009) Dynamic pricing and advertising for web content providers. *European Journal of Operational Research* 197(3):924–944.
- Lambrecht A, Tucker C (2013) When does retargeting work? information specificity in online advertising. *Journal of Marketing Research* 50(5):561–576.
- Lambrecht A, Tucker CE (2017) Algorithmic bias? an empirical study into apparent gender-based discrimination in the display of stem career ads Working paper.
- Lang KJ, Moseley B, Vassilvitskii S (2012) Handling forecast errors while bidding for display advertising. *Proceedings of the 21st international conference on world wide web*, 371–380 (ACM).
- Lee D, Hosanagar K, Nair H (2017a) Advertising content and consumer engagement on social media: Evidence from facebook Working paper.
- Lee GM, He S, Lee J, Whinston AB (2017b) Matching mobile applications for cross promotion .
- Lee KC, Jalali A, Dasdan A (2013) Real time bid optimization with smooth budget delivery in online advertising. *Proceedings of the 7th international workshop on data mining for online advertising*, 1 (ACM).
- Levin J, Milgrom P (2010) Online advertising: Heterogeneity and conflation in market design. *The American Economic Review* 100(2):603–607.
- Lewis R, Nguyen D (2015) Display advertising’s competitive spillovers to consumer search. *Quantitative Marketing and Economics* 13(2):93–115.
- Lewis R, Rao JM, Reiley DH (2015) Measuring the effects of advertising: The digital frontier. *Economic Analysis of the Digital Economy*, 191–218 (University of Chicago Press).
- Lewis RA, Reiley DH (2014) Online ads and offline sales: Measuring the effect of retail advertising via a controlled experiment on yahoo! *Quantitative Marketing and Economics* 12(3):235–266.
- Lewis RA, Wong J (2018) Incrementality bidding & attribution .
- Li C, Luo X, Zhang C, Wang X (2017) Sunny, rainy, and cloudy with a chance of mobile promotion effectiveness. *Marketing Science* 36(5):762–779.
- Li H, Kannan P (2014) Attributing conversions in a multichannel online marketing environment: An empirical model and a field experiment. *Journal of Marketing Research* 51(1):40–56.
- Liu D, Kumar S, Mookerjee VS (2012) Advertising strategies in electronic retailing: A differential games approach. *Information Systems Research* 23(3-part-2):903–917.
- Liu D, Viswanathan S (2013) Information asymmetry and hybrid advertising. *Journal of Marketing Research* 50(5):609–624.
- Liu K, Tang L (2011) Large-scale behavioral targeting with a social twist. *Proceedings of the 20th international conference on information and knowledge management*, 1815–1824 (ACM).
- Lu S, Yang S (2015) A two-sided market analysis of behaviorally targeted display advertising Working paper.
- Luo X, Andrews M, Fang Z, Phang CW (2013) Mobile targeting. *Management Science* 60(7):1738–1756.
- Manchanda P, Dubé JP, Goh KY, Chintagunta PK (2006) The effect of banner advertising on internet purchasing. *Journal of Marketing Research* 43(1):98–108.
- Mangani A (2004) Online advertising: Pay-per-view versus pay-per-click. *Journal of Revenue and Pricing Management* 2(4):295–302.
- Marotta V, Zhang K, Acquisti A (2017) The welfare impact of targeted advertising Working paper.
- McMahan HB, Holt G, Sculley D, Young M, Ebner D, Grady J, Nie L, Phillips T, Davydov E, Golovin D, et al. (2013) Ad click prediction: A view from the trenches. *Proceedings of the 19th international conference on knowledge discovery and data mining*, 1222–1230 (ACM).
- Medina AM, Mohri M (2014) Learning theory and algorithms for revenue optimization in second price auctions with reserve. *Proceedings of the 31st international conference on machine learning*, 262–270.
- Menon AK, Chitrapura KP, Garg S, Agarwal D, Kota N (2011) Response prediction using collaborative filtering with hierarchies and side-information. *Proceedings of the 17th international*

- conference on knowledge discovery and data mining*, 141–149 (ACM).
- Mirrokní V, Nazerzadeh H (2017) Deals or no deals: Contract design for online advertising. *Proceedings of the 26th international conference on world wide web*, 7–14 (International World Wide Web Conferences Steering Committee).
- Mirrokní V, Paes Leme R, Tang P, Zuo S (2018) Non-clairvoyant dynamic mechanism design. *Available at SSRN 2873701*.
- Mookerjee R, Kumar S, Mookerjee VS (2016) Optimizing performance-based internet advertisement campaigns. *Operations Research* 65(1):38–54.
- Muthukrishnan S (2009) Ad exchanges: Research issues. *International workshop on internet and network economics*, 1–12 (Springer).
- Myerson RB (1981) Optimal auction design. *Mathematics of Operations Research* 6(1):58–73.
- Najafi-Asadolahi S, Fridgeirsdottir K (2014) Cost-per-click pricing for display advertising. *Manufacturing & Service Operations Management* 16(4):482–497.
- Oentaryo RJ, Lim EP, Finegold M, Lo D, Zhu F, Phua C, Cheu EY, Yap GE, Sim K, Nguyen MN, et al. (2014a) Detecting click fraud in online advertising: A data mining approach. *Journal of Machine Learning Research* 15(1):99–140.
- Oentaryo RJ, Lim EP, Low JW, Lo D, Finegold M (2014b) Predicting response in mobile advertising with hierarchical importance-aware factorization machine. *Proceedings of the 7th international conference on web search and data mining*, 123–132 (ACM).
- Paes Leme R, Pál M, Vassilvitskii S (2016) A field guide to personalized reserve prices. *Proceedings of the 25th international conference on world wide web*, 1093–1102 (International World Wide Web Conferences Steering Committee).
- Pandey S, Aly M, Bagherjeiran A, Hatch A, Ciccolo P, Ratnaparkhi A, Zinkevich M (2011) Learning to target: What works for behavioral targeting. *Proceedings of the 20th international conference on information and knowledge management*, 1805–1814 (ACM).
- Papadimitriou P, Garcia-Molina H, Krishnamurthy P, Lewis RA, Reiley DH (2011) Display advertising impact: Search lift and social influence. *Proceedings of the 17th international conference on knowledge discovery and data mining*, 1019–1027 (ACM).
- Paulson C, Luo L, James GM (2018) Efficient large-scale internet media selection optimization for online display advertising. *Journal of Marketing Research* 55(4):489–506.
- Perlich C, Dalessandro B, Hook R, Stitelman O, Raeder T, Provost F (2012) Bid optimizing and inventory scoring in targeted online advertising. *Proceedings of the 18th international conference on knowledge discovery and data mining*, 804–812 (ACM).
- Perlich C, Dalessandro B, Raeder T, Stitelman O, Provost F (2014) Machine learning for targeted display advertising: Transfer learning in action. *Machine learning* 95(1):103–127.
- Provost F, Dalessandro B, Hook R, Zhang X, Murray A (2009) Audience selection for on-line brand advertising: privacy-friendly social network targeting. *Proceedings of the 15th international conference on knowledge discovery and data mining*, 707–716 (ACM).
- Provost F, Martens D, Murray A (2015) Finding similar mobile consumers with a privacy-friendly geosocial design. *Information Systems Research* 26(2):243–265.
- Rafieian O, Yoganarasimhan H (2018a) How does variety of previous ads influence consumer’s ad response? Working paper.
- Rafieian O, Yoganarasimhan H (2018b) The value of information in mobile ad targeting. *Available at SSRN 3163806* Working paper.
- Ren K, Zhang W, Rong Y, Zhang H, Yu Y, Wang J (2016) User response learning for directly optimizing campaign performance in display advertising. *Proceedings of the 25th international conference on information and knowledge management*, 679–688 (ACM).
- Roels G, Fridgeirsdottir K (2009) Dynamic revenue management for online display advertising. *Journal of Revenue and Pricing Management* 8(5):452–466.
- Sahni NS (2015) Effect of temporal spacing between advertising exposures: Evidence from online field experiments. *Quantitative Marketing and Economics* 13(3):203–247.
- Sahni NS, Narayanan S, Kalyanam K (2019) An experimental investigation of the effects of retargeted advertising: The role of frequency and timing. *Journal of Marketing Research*.
- Salz T (2017) Intermediation and competition in search markets: An empirical case study Working paper.
- Sayedí A (2018) Real-time bidding in online display advertising. *Marketing Science* 37(4):553–568.

- Sayed A, Jerath K, Baghaie M (2018) Exclusive placement in online advertising. *Marketing Science* 37(6):970–986.
- Schwartz EM, Bradlow ET, Fader PS (2017) Customer acquisition via display advertising using multi-armed bandit experiments. *Marketing Science* 36(4):500–522.
- Scott SL (2010) A modern bayesian look at the multi-armed bandit. *Applied Stochastic Models in Business and Industry* 26(6):639–658.
- Shao X, Li L (2011) Data-driven multi-touch attribution models. *Proceedings of the 17th international conference on knowledge discovery and data mining*, 258–264 (ACM).
- Stitelman O, Perlich C, Dalessandro B, Hook R, Raeder T, Provost F (2013) Using co-visitation networks for detecting large scale online display advertising exchange fraud. *Proceedings of the 19th international conference on knowledge discovery and data mining*, 1240–1248 (ACM).
- Stone-Gross B, Stevens R, Zarras A, Kemmerer R, Kruegel C, Vigna G (2011) Understanding fraudulent activities in online ad exchanges. *Proceedings of the conference on internet measurement conference*, 279–294 (ACM).
- Stourm V, Bax E (2017) Incorporating hidden costs of annoying ads in display auctions. *International Journal of Research in Marketing* 34(3):622–640.
- Sun Z, Dawande M, Janakiraman G, Mookerjee V (2017) Not just a fad: Optimal sequencing in mobile in-app advertising. *Information Systems Research* 28(3):511–528.
- Sun Z, Dawande M, Janakiraman G, Mookerjee VS (2016) The making of a good impression: Information hiding in ad exchanges. *MIS Quarterly* 40(3):717–739.
- Sweeney L (2013) Discrimination in online ad delivery. *Queue* 11(3):10.
- Toubiana V, Narayanan A, Boneh D, Nissenbaum H, Barocas S (2010) Adnostic: Privacy preserving targeted advertising. *Proceedings on network and distributed system symposium*.
- Trusov M, Ma L, Jamal Z (2016) Crumbs of the cookie: User profiling in customer-base analysis and behavioral targeting. *Marketing Science* 35(3):405–426.
- Tucker C (2012a) The economics of advertising and privacy. *International Journal of Industrial Organization* 30(3):326–329.
- Tucker C (2012b) Social advertising Working paper.
- Tucker C (2014a) The reach and persuasiveness of viral video ads. *Marketing Science* 34(2):281–296.
- Tucker C (2014b) Social networks, personalized advertising, and privacy controls. *Journal of Marketing Research* 51(5):546–562.
- Tunuguntla S, Hoban PR (2018) A near optimal bidding strategy for real-time display advertising auctions Working paper.
- Turner J (2012) The planning of guaranteed targeted display advertising. *Operations Research* 60(1):18–33.
- Urban GL, Liberali G, MacDonald E, Bordley R, Hauser JR (2013) Morphing banner advertising. *Marketing Science* 33(1):27–46.
- Vee E, Vassilvitskii S, Shanmugasundaram J (2010) Optimal online assignment with forecasts. *Proceedings of the 11th conference on electronic commerce*, 109–118 (ACM).
- Wang J, Zhang W (2016) Bid-aware gradient descent for unbiased learning with censored data in display advertising. *Proceedings of the 22nd international conference on knowledge discovery and data mining*, volume 22, 665–674 (ACM).
- Wang J, Zhang W, Yuan S (2016) Display advertising with real-time bidding (rtb) and behavioural targeting Working paper.
- Wojdyski BW, Evans NJ (2016) Going native: Effects of disclosure position and language on the recognition and evaluation of online native advertising. *Journal of Advertising* 45(2):157–168.
- Wu WCH, Yeh MY, Chen MS (2015) Predicting winning price in real time bidding with censored data. *Proceedings of the 21th international conference on knowledge discovery and data mining*, 1305–1314 (ACM).
- Xu J, Lee Kc, Li W, Qi H, Lu Q (2015) Smart pacing for effective online ad campaign optimization. *Proceedings of the 21th international conference on knowledge discovery and data mining*, 2217–2226 (ACM).
- Xu J, Shao X, Ma J, Lee Kc, Qi H, Lu Q (2016) Lift-based bidding in ad selection Working paper.
- Xu L, Duan JA, Whinston A (2014) Path to purchase: A mutually exciting point process model for online advertising and conversion. *Management Science* 60(6):1392–1412.

- Yuan S, Wang J, Chen B, Mason P, Seljan S (2014) An empirical study of reserve price optimisation in real-time bidding. *Proceedings of the 20th international conference on knowledge discovery and data mining*, 1897–1906 (ACM).
- Yuan S, Wang J, Zhao X (2013) Real-time bidding for online advertising: Measurement and analysis. *Proceedings of the 7th international workshop on data mining for online advertising*, 3 (ACM).
- Zeithammer R (2016) The futility of soft floor auctions .
- Zhang K, Bhattacharyya S, Ram S (2016a) Large-scale network analysis for online social brand advertising. *Mis Quarterly* 40(4).
- Zhang K, Katona Z (2012) Contextual advertising. *Marketing Science* 31(6):980–994.
- Zhang W, Du T, Wang J (2016b) Deep learning over multi-field categorical data. *European Conference on Information Retrieval*, 45–57 (Springer).
- Zhang W, Yuan S, Wang J (2014) Optimal real-time bidding for display advertising. *Proceedings of the 20th international conference on knowledge discovery and data mining*, 1077–1086 (ACM).
- Zhu Y, Wilbur KC (2011) Hybrid advertising auctions. *Marketing Science* 30(2):249–273.
- Zubcsek PP, Sarvary M (2011) Advertising to a social network. *Quantitative Marketing and Economics* 9(1):71–107.

Online Appendix

A Video

Two-thirds of advertisers surveyed by IAB reported they will shift funds from TV to digital video ads.⁴⁵ Yet little is known about the role of video ads related to advertiser objectives and how they interact with other formats/channels such as banner ads or traditional TV ads. At a finer level, questions exist regarding (i) which forms (e.g. in-stream, out-stream, in-display), (ii) which websites (e.g. social media), and (iii) which devices moderate the efficacy of video ads. For in-stream video ads, the timing (pre-, mid-, post-roll), the sequence of competing advertisers, duration, and skippability (Chiong et al. 2018) may also affect ad effectiveness in a similar or dissimilar manner to TV advertising.

At a higher display ecosystem level, video ads require fast and high quality rendering, and the complexity of integration rises further as publishers consider selling video inventory in real-time. Liu et al. (2012) analyze a situation where advertising expenditure can become wasteful under an IT capacity constraint, as consumers may be less likely to buy a product from the focal advertiser’s website due to, for example, slow real-time targeted video rendering.⁴⁶ This issue suggests that investment on technology becomes an additional consideration factor in buying, matching, and selling video ads. Finally, as video ads are a relative nascent format, standardizing the measurement results and developing comparable metrics across platforms (e.g., TV and digital video) are becoming priorities for both practitioners and researchers alike.⁴⁷

B Transparency

There is an increasing demand for transparency in the display ad market, more so with the complexity arising in RTB and with the growth of fraudulent activities (Daswani et al. 2008; Feily et al. 2009; Stone-Gross et al. 2011). In this section, we discuss the industry’s effort to adopt better measures (e.g., eligible, viewable, rendered) and research issues related to transparency.

Not all impressions bought are actually viewed by the users, leading to difficulties in predicting ROI and making ad buying decisions at appropriate prices.⁴⁸ More advertisers are requesting

⁴⁵ https://www.iab.com/wp-content/uploads/2018/04/2018_IAB_NewFronts_Video_Ad_Spend_Report.pdf

⁴⁶ <https://video-guide.iab.com/challenges-digital-video-space-and-efforts-underway-solve-them>

⁴⁷ <https://video-guide.iab.com/challenges-digital-video-space-and-efforts-underway-solve-them>

⁴⁸ The traditional metric, impression, is counted when an ad creative is sent by the ad server. The counted impression may not be actually viewed by the user, for example, when he/she does not scroll down far enough to see ads loaded at

performance reports with metrics, such as *viewable* impressions, and are demanding pricing schemes to reflect these alternative metrics (e.g., paying only for viewable impressions).⁴⁹ In order to provide assurance to advertisers, publishers often use third parties (e.g., eTrust) to verify that they don't serve ads to impressions generated by botnets.⁵⁰ Yet, many publishers are wary of meeting the viewability standards and of the prospect of losing revenues when incorporating the viewability metric into pricing schemes.⁵¹ Understanding the impact of viewability metrics on the profits of advertisers, publishers, and intermediaries is an imminent research question. There exists a stream of research that develops algorithms for detecting fraud impressions and clicks (Stitelman et al. 2013; Crussell et al. 2014; Oentaryo et al. 2014a), and an area of future research is to examine the players' (publishers' or intermediaries') incentives to adopt these algorithms (Chen et al. 2015 under CPC).

Last, there exist concerns that ad targeting/serving algorithms may result in discriminatory outcomes, intentionally or unintentionally, leading to calls for algorithmic transparency (Datta et al. 2015; Sweeney 2013). For example, Lambrecht & Tucker (2017) find competitive spillover in STEM career ads, where female impressions are highly valued by competing advertisers. This leads to displaying more STEM ads to men by bidding algorithms designed to be cost-effective (and intended to be gender-neutral). Delving further into the underlying mechanisms across various contexts could yield consequential policy insights toward preventing undesirable biases.

C A Brief History of Display Advertising

Display advertising began in 1994 when HotWired invented and sold the first web banner ad space to AT&T at \$30,000 for running a three-month campaign. This ad, with the copy "Have you ever clicked your mouse right here? You will", generated a 44% CTR.⁵² This figure is astonishing when compared to today's average rate of 0.05%.⁵³

the bottom of a page, or when he/she uses ad blocking tools to purposefully avoid ads. In addition, fake impressions can be generated by fraudulent activities involving botnets or pop-under windows with 0x0 size pixels.

⁴⁹ According to the IAB's guideline, an impression is considered viewable when at least 50% of an ad's pixels are in view on the in-focus browser for a minimum of one continuous second (http://mediaratingcouncil.org/081815%20Viewable%20Ad%20Impression%20Guideline_v2.0_Final.pdf).

⁵⁰ Third parties, for example, use blockchain to audit ad transactions (Catalini & Gans 2016).

⁵¹ <http://blog.imonomy.com/meeting-the-viewability-threat-to-publishers-revenue/>
<https://digiday.com/media/publishers-still-bristle-groupms-tougher-viewability-standards/>

⁵² <https://blog.hubspot.com/marketing/history-of-online-advertising>

⁵³ <https://www.smartinsights.com/internet-advertising/internet-advertising-analytics/display-advertising-clickthrough-rates/>

With the increase in popularity for banner ads came the emergence of ad tech companies who created new ways to help advertisers better manage and measure the results of ad campaigns. A notable entrant is DoubleClick who offered tools to track views/clicks across multiple websites on its network and gave rise to the CPM pricing model (instead of a flat fee for a campaign).⁵⁴ In spite of its early promise, display advertising gradually yielded its advertising market share to sponsored search advertising in the 2000s as search engines became gateways for consumers navigating through a myriad of websites and as advertisers looked for better ways to target users based on their keyword inputs. By the end of the 2000s, sponsored search (display) advertising attained 46% (38%) of online advertising spend, as compared to 1% (78%) in 2000.⁵⁵

More recently, display advertising has regained its momentum in the 2010s with the rapid increase in social media activities, mobile browsing, and video ad formats.⁵⁶ Both advertisers and publishers are moving toward enhancing users' online experiences by increasing quality of the ad content, format, and targeting.⁵⁷ Less intrusive native advertising, which matches the content and style of the ad creative to the form and function of the website, has been increasing in popularity. The resurgence of display advertising has also been fueled by the concurrent advancement in enabling technologies (e.g., RTB), and this market is expected to continue its growth affording numerous opportunities for researchers as discussed in the paper.

D Display vs. Search Advertising Markets

Although some aspects of display advertising are similar to sponsored search advertising (e.g., auction mechanism is used for selling ads), there exist unique characteristics that distinguish display ads from search ads. On the advertiser side, one difference between these two advertising types pertains to the customer's purchase intent: search advertising by definition requires keyword inputs from the users, whereas display advertising is delivered to those who are browsing the web. As a result, search advertising is more likely to reach those at the point of purchase. Display advertising, on the other hand, can be targeted at the higher (e.g., brand building) and lower (e.g., retargeting) levels of the purchase funnel. Another distinctive decision for the advertisers is creative. Unlike

⁵⁴ <https://www.adpushup.com/blog/the-history-of-online-advertising/>

⁵⁵ www.iab.com

⁵⁶ <https://www.iab.com/wp-content/uploads/2018/05/IAB-2017-Full-Year-Internet-Advertising-Revenue-Webinar-Presentation.pdf>

⁵⁷ <https://www.cnn.com/2015/04/14/facebook-stands-behind-its-controversial-ad-strategy.html>

search ads, display advertising delivers powerful visual effects that can enhance engagement and user experience.

From the seller’s point of view, there are two main differences worth mentioning. One is that the publisher in search advertising (i.e., search engine platform) is vertically integrated, leading to a much simpler market structure with few or no intermediaries. Similar to the vertically integrated Facebook, search engine platforms face the trade-off between showing organic results of high relevance and ads yielding high revenue. This trade-off is less salient in (traditional websites) display advertising, where the publisher’s decision is to assign a portion of advertising space in which ads do not compete directly with the product/content/service provided. Another important distinction arises from the unit of sale and the associated auction mechanism. In major search engine platforms, each keyword is sold based on the cost-per-click (CPC) pricing model via generalized second-price auction thereby triggering the simultaneous sale of multiple ad slots (in contrast to display markets). As such, advertisers in search ad markets compete for better position ranks for a given keyword. To balance users’ experiences and advertising revenues, search engine platforms often incorporate an ad relevance score (together with bid) in determining the winner of the auction. In contrast, the unit of sale in display advertising is each user’s impression for a given placement, and the commonly adopted auction format is second-price, sealed-bid auction based on cost-per-mille (CPM) (as opposed to a position auction in search). Lastly, search queries are sold on the real-time basis and advertisers are not guaranteed in advance to receive a certain number of impressions (or clicks). In display advertising, both guaranteed and non-guaranteed markets coexist.

Overall, the difference in advertiser goals, ad formats, auction mechanisms, intermediaries, inventory, channel structure, and purchase timing imply there exist some considerable differences between display and search ad markets.